



Datomic

@stuarthalloway

the Datomic database is

indelible

chronological

flexible

powerful

simple

Datomic is *indelible*:

datoms cannot be
modified or removed

why source control?

reify change

trigger workflows

compare points in time

prove a point (audit)

generic (vs. ad hoc) model of time

update-in-place (sad)

entity	attribute	value
jane	likes	broccoli

update-in-place (sad)

entity	attribute	value
jane	likes	pizza

indelible, not CRUD

“create”

assertion

“update”

retraction + assertion

“delete”

retraction

indelible

entity	attribute	value	transaction	assert?
jane	likes	broccoli	1008	true
jane	likes	broccoli	1148	false
jane	likes	pizza	1148	true

Datomic is chronological

every datom is timestamped

query as-of a particular moment in time

strong consistency

ACID

CAP

filtered retractions

entity	attribute	value	transaction	assert?
jane	likes	broccoli	1008	true
1008	txInstant	... 04:00	1008	true
jane	likes	broccoli	1148	false
jane	likes	pizza	1148	true
1148	txInstant	... 03:00	1148	true

as-of

entity	attribute	value	transaction	assert?
jane	likes	broccoli	1008	true
1008	txInstant	... 04:00	1008	true
jane	likes	broccoli	1148	false
jane	likes	pizza	1148	true
1148	txInstant	... 03:00	1148	true

history

entity	attribute	value	transaction	assert?
jane	likes	broccoli	1008	true
1008	txInstant	... 04:00	1008	true
jane	likes	broccoli	1148	false
jane	likes	pizza	1148	true
1148	txInstant	... 03:00	1148	true

“Code doesn’t exist
unless it’s checked into a
version control system”

“Data doesn’t exist unless
it’s transacted into an
indelible, chronological
database”

Datomic is flexible

Datomic stores granular *information*

datoms are the atoms of information

attribute-level schema

model your domain

instead of torturing it to fit into tables

trading off flexibility (NoSQL)

“picking the right data model is the hardest part ...”

“model your data to fit your queries”

“don’t model around relations”

“don’t model around objects”

flexibility

picking the right data model is the easiest part

model your data to fit your domain

model relations

model objects

no “impedance mismatch”

one db, many query styles

structure	attribute
k/v	AVET
row	EAVT
column	AEVT
document	EAVT, components
graph	VAET

Datomic is powerful

datalog query

logic programming

pattern syntax

joins

rules

example database

entity	attribute	value
42	:email	<u>jdoo@exampl.com</u>
43	:email	<u>jane@exampl.com</u>
42	:orders	107
42	:orders	141

data pattern

*Constrains the results returned,
binds variables*

variable



variable



[?customer :email ?email]

find by email

entity	attribute	value
42	:email	<u>jdoe@example.com</u>
43	:email	<u>jane@example.com</u>
42	:orders	107
42	:orders	141

[?customer :email ?email]

the gourmet jerky
problem

```
[ :find ?customer ?product
  :where [?customer :shipAddress ?addr]
          [?addr :zip ?zip]
          [?product :product/weight ?weight]
          [?product :product/price ?price]
          [(Shipping/estimate ?zip ?weight) ?shipCost]
          [(<= ?price ?shipCost)]]
```


Datomic is simple

simple

not complected

not woven together

orthogonal

a complected database

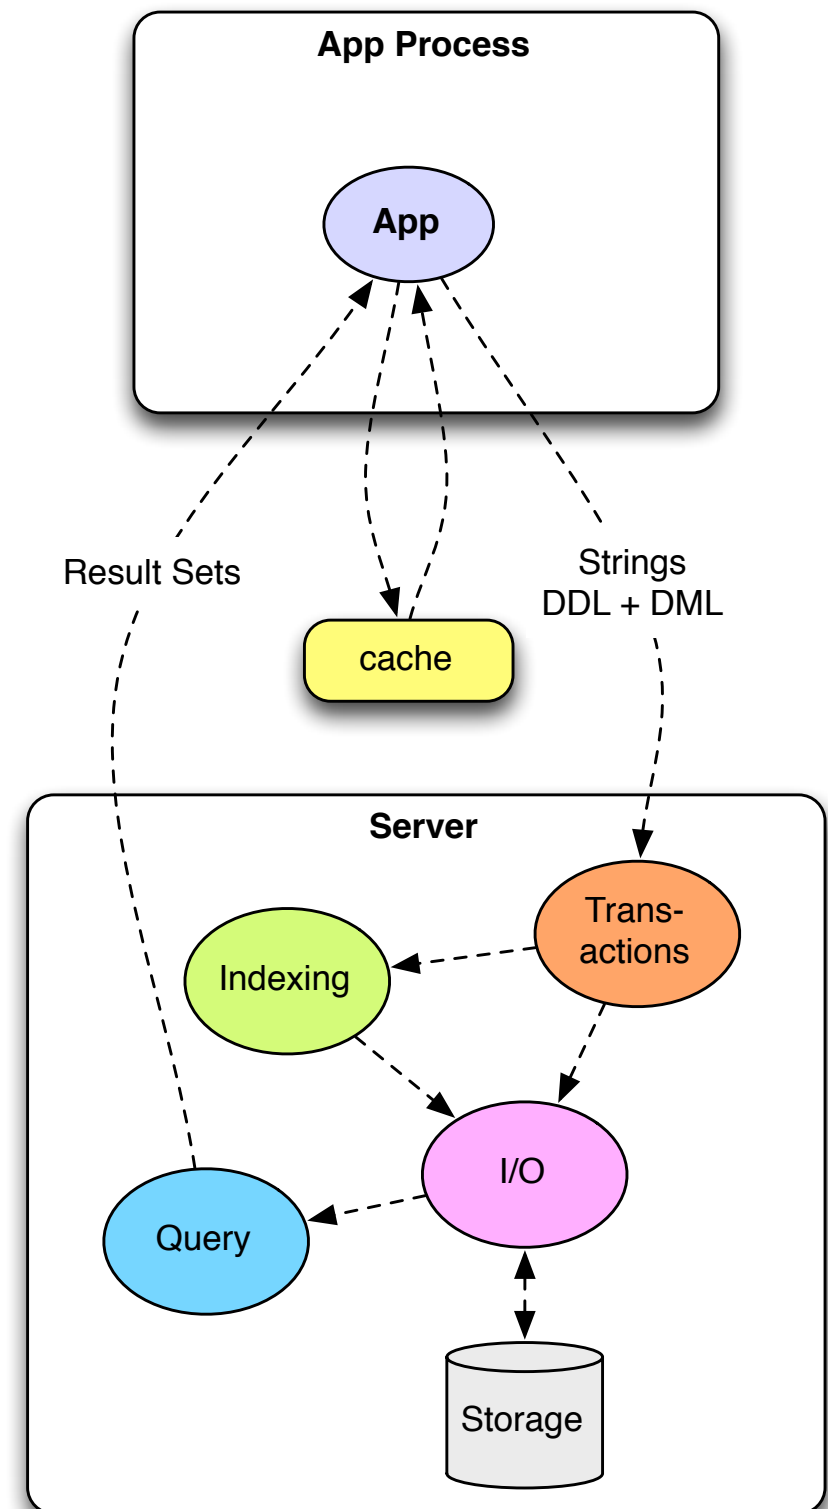
monolithic server

reads


writes

indexing

storage



Datomic simplicity



architectural benefits

run on your own storage

DynamoDB

SQL

Cassandra

horizontal read scaling

Datomic is

indelible

chronological

flexible

powerful

simple

successful

“We needed the flexibility and agility of a startup, and the system of record / audit trail features one would find in legacy banking systems. I didn’t find a better option than **Datomic with this first class concept of time**”

how did a small team
build this?



what made lisp different

feature	Java	Clojure
conditionals	✓	✓
variables	✓	✓
garbage collection	✓	✓
recursion	✓	✓
function type	*	✓
symbol type		✓
whole language available		✓
everything's an expression		✓
homoiconicity		✓

ask the internet

rare

cuts through any problem

intimidating

elegant

ask the internet

rare



cuts through any problem

intimidating

elegant

**instantly recognizable emblem of people you don't want
to f*** with**

<https://xkcd.com/297/>



extensible data notation (edn)

```
{ :firstName "John"  
  :lastName "Smith"  
  :age 25  
  :address {  
    :streetAddress "21 2nd Street"  
    :city "New York"  
    :state "NY"  
    :postalCode "10021" }  
  :phoneNumber  
    [ { :type "name" :number "212 555-1234"}  
      { :type "fax" :number "646 555-4567" } ] }
```

type	examples
string	<code>"foo"</code>
character	<code>\f</code>
integer	<code>42, 42N</code>
floating point	<code>3.14, 3.14M</code>
boolean	<code>true</code>
nil	<code>nil</code>
symbol	<code>foo, +</code>
keyword	<code>:foo, ::foo</code>

type	properties	examples
list	sequential	(1 2 3)
vector	sequential and random access	[1 2 3]
map	associative	{ :a 100 :b 90 }
set	membership	# { :a :b }

"hello world"

hello you

define a fn fn name docstring

```
(defn greet  
  "Returns a friendly greeting"  
  [you]  
  (str "Hello, " you))
```

arguments fn body

the StrUtils problem

existing class

existing interface

never the twain shall meet

protocols

```
(defprotocol Blank  
  (blank? [_]))
```


```
(blank? "foo")  
=> IllegalArgumentException
```

```
(extend-protocol Blank  
  String  
  (blank?  
    [s]  
    (every? #(Character/isspace %) s)))
```

```
(blank? "  ")
```

```
=> true
```

storage protocol



testing is hard

difficult to interpret tests as knowledge about system

difficult to achieve good coverage

costly to develop and maintain

trivial and serial tests vs. complex and parallel reality

generative testing

programmer models the domain

a *program* writes the individual tests

validate categoric properties, not specific outcomes

a program wrote this program

```
;;;;;;;;;;;;;;  
;; Case 0  
  
(def  
  sql0  
  "SELECT r1.e,r1.v\nFROM r1\nWHERE (r1.e=0 OR r1.e=1 OR r1.e=7) AND r1.a='b'")  
(def e0 "Expected result for test query 0" #{[7 1] [0 8]})  
(let [expected e0 actual  
      (datomic.api/q  
        (quote  
          {:find [?r1-e ?r1-v],  
            :where ([$r1 ?r1-e :b ?r1-v]  
                    [(contains? #{0 7 1} ?r1-e))],  
            :in [$r1]) db1)]  
      (when-not (= expected actual)  
        (throw (ex-info "Query results did not crosscheck"  
                        { :expected expected, :actual actual}))))
```

reading test failures is hard

```
(tc/quick-check 100 some-complex-property)
```

```
{:fail [[10 1 28 40 11 -33 42 -42 39 -13  
          13 -44 -36 11 27 -42 4 21 -39]]}
```


programmatic shrinking

```
(tc/quick-check 100 some-complex-property)
```

```
{:fail [[10 1 28 40 11 -33 42 -42 39 -13  
        13 -44 -36 11 27 -42 4 21 -39]],  
 :shrunk {:total-nodes-visited 38,  
          :depth 18,  
          :result false,  
          :smallest [[42]]}}
```

communication is hard

objects make terrible machines

function chains make poor machines

direct-connect relationships

callback hell

j.u.c queues block real threads

threads are expensive and/or nonexistent

core.async (CSP)

first class processes

first class channels

concurrency primitive (coordination)

coherent sequential logic

multi reader/writer

buffering

search with no threads, SLA

```
(defn search [query]
  (let [c (chan)
        t (timeout 80)]
    (go (>! c (<! (fastest query web1 web2))))
    (go (>! c (<! (fastest query image1 image2))))
    (go (>! c (<! (fastest query video1 video2))))
    (go (loop [i 0
              ret []]
        (if (= i 3)
          ret
          (recur (inc i)
                 (conj ret (alt! [c t] ([v] v))))))))))
```

basic transduction

```
(a/transduce
  (comp
    (halt-when error?)
    (map ...)
    (filter ...)))
  (completing ...)
  accumulator
  query)
```

clojure.spec

a standard, expressive,
powerful, integrated system for
specification and testing

a taste of spec

```
(s/def ::sku  
  (s/and string?  
    #(str/starts-with "SKU-" %)))
```

```
(s/def ::purchaser  
  (s/or :account-id pos-int?  
        :email string?))
```

```
(s/def ::quantity pos-int?)
```

```
(s/def ::import-record  
  (s/tuple [::purchaser ::sku ::quantity]))
```

spec capabilities

what	how
what are the building blocks?	declarative structure
what invariants hold?	arbitrary predicates
how do I check?	validation
what went wrong?	explanation
what went right?	conformance
docs	autodoc
examples	sample generator
am I using this right?	instrumentation
is my code correct?	generative testing
a la carte self-check	assertion

spec vs. Java types and tests

	jUnit etc.	Java types	spec
expressive	very	no	very
powerful	stakeholder correctness	type correctness	stakeholder correctness
integrated	no	compile-time, must flow	dynamic
specification	no	static	yes
testing	manual	no	generative
agility	expensive	fragility	dynamic
reach	expensive	libs, apps	systemic

core values

simplicity

power

focus

stability

simplicity

just keep finding smaller abstractions?!

planning, agility, and patience

takes time and courage

power

design for maximum leverage from the platform (JVM)

semantic fidelity

performance

semantic fidelity


Clojure programs are Java programs

call Java from Clojure without wrappers

call Clojure from Java without wrappers

performance

note:
log
scale!



<http://benchmarkgame.alioth.debian.org/u64q/which-programs-are-fastest.php>

focus

related coding principles

- single-responsibility principle

- don't repeat yourself (dry)

Clojure superpowers

- simplicity & homoiconicity

- systemic generality

systemic generality

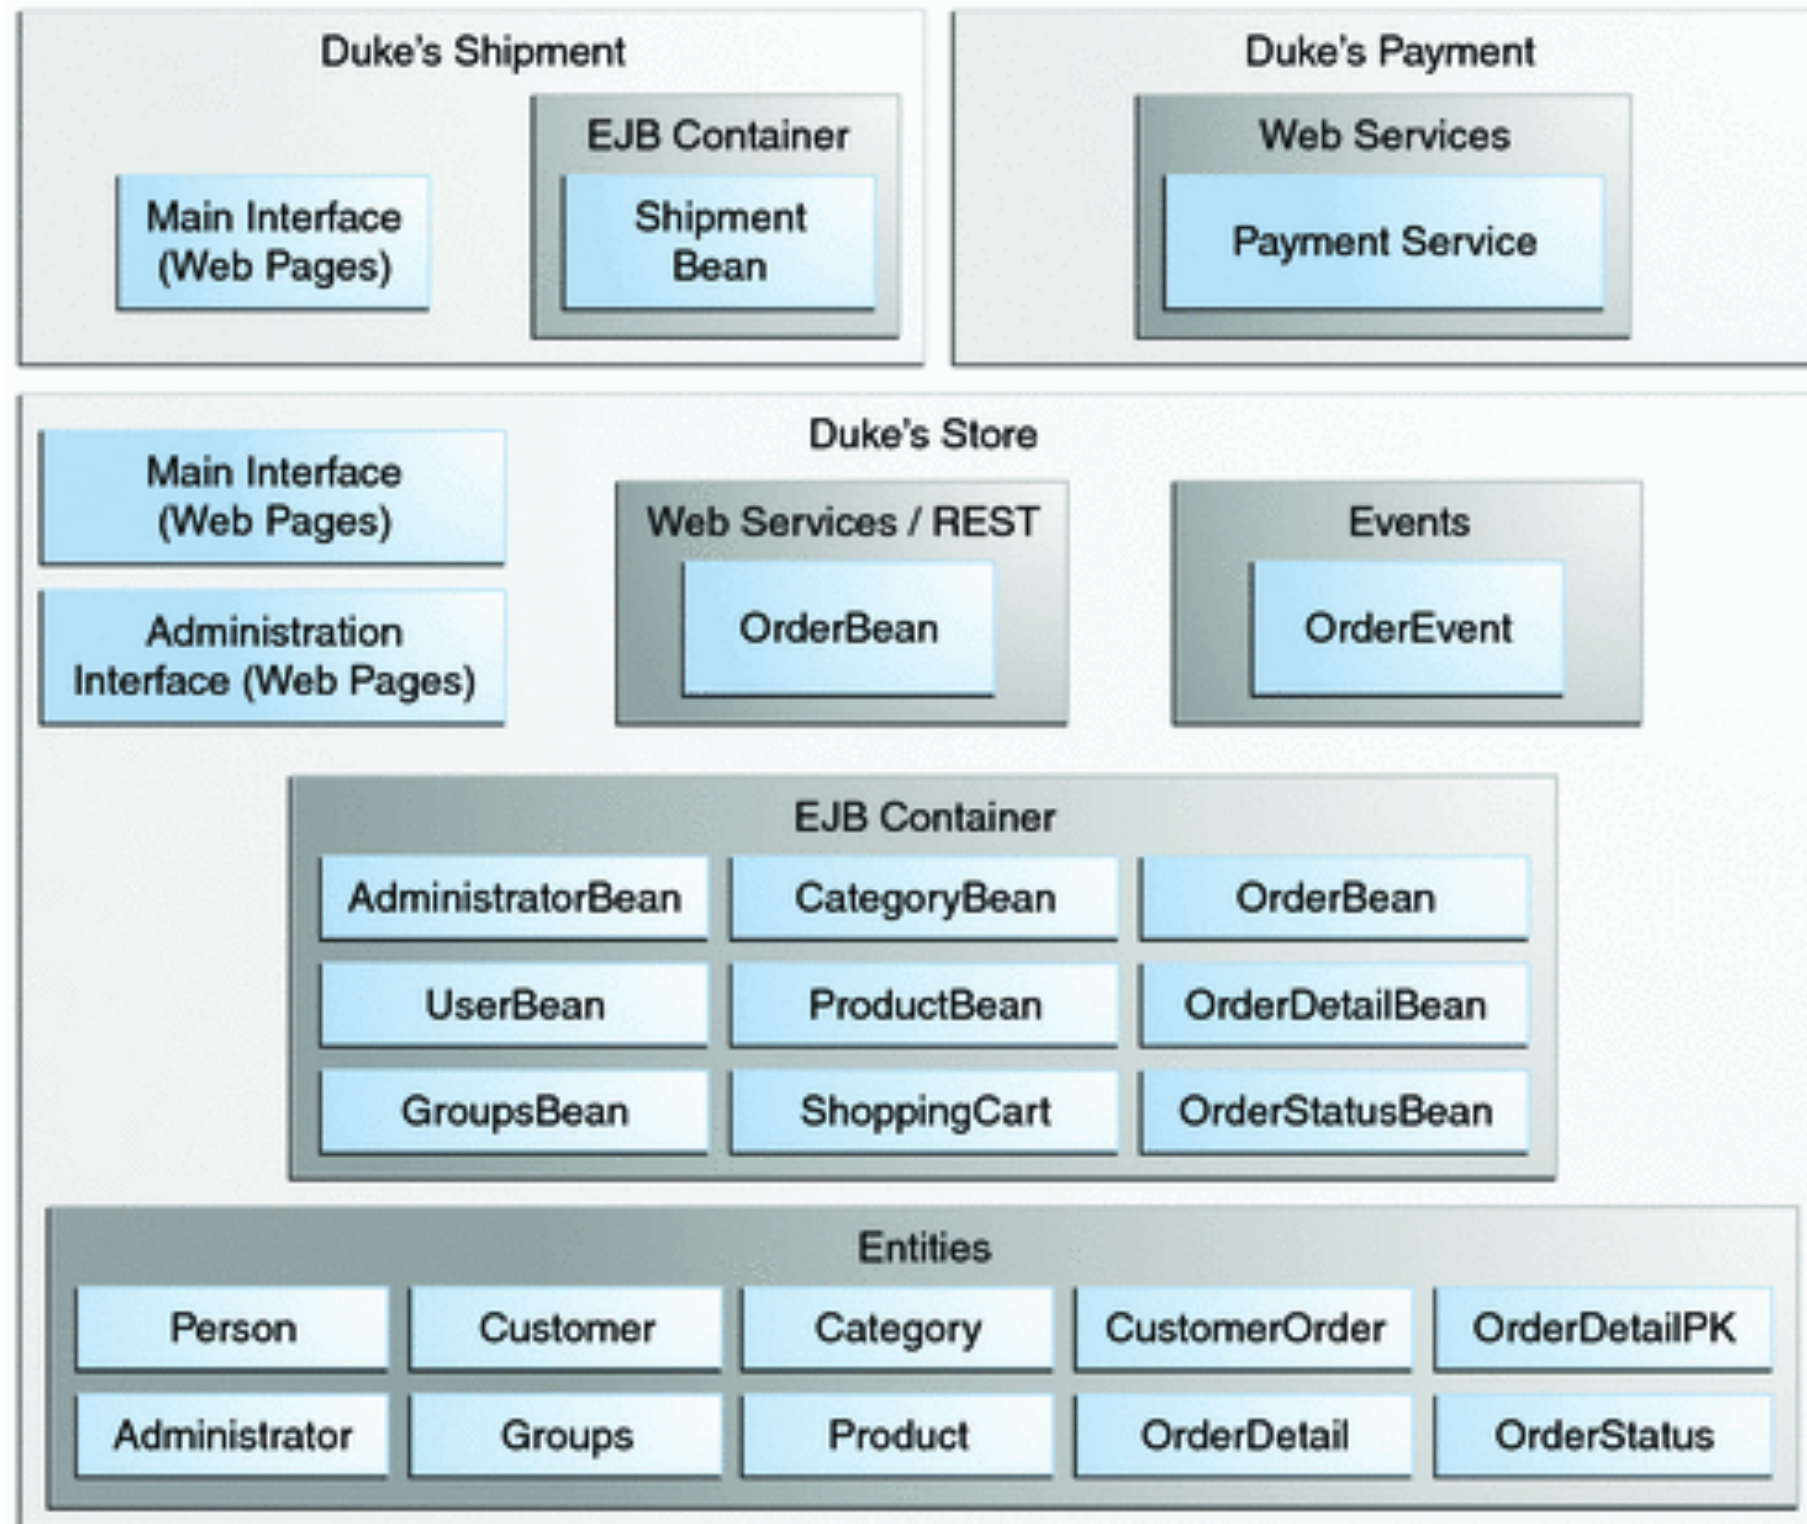
generality

all domains use the same general-purpose
data structures and functions

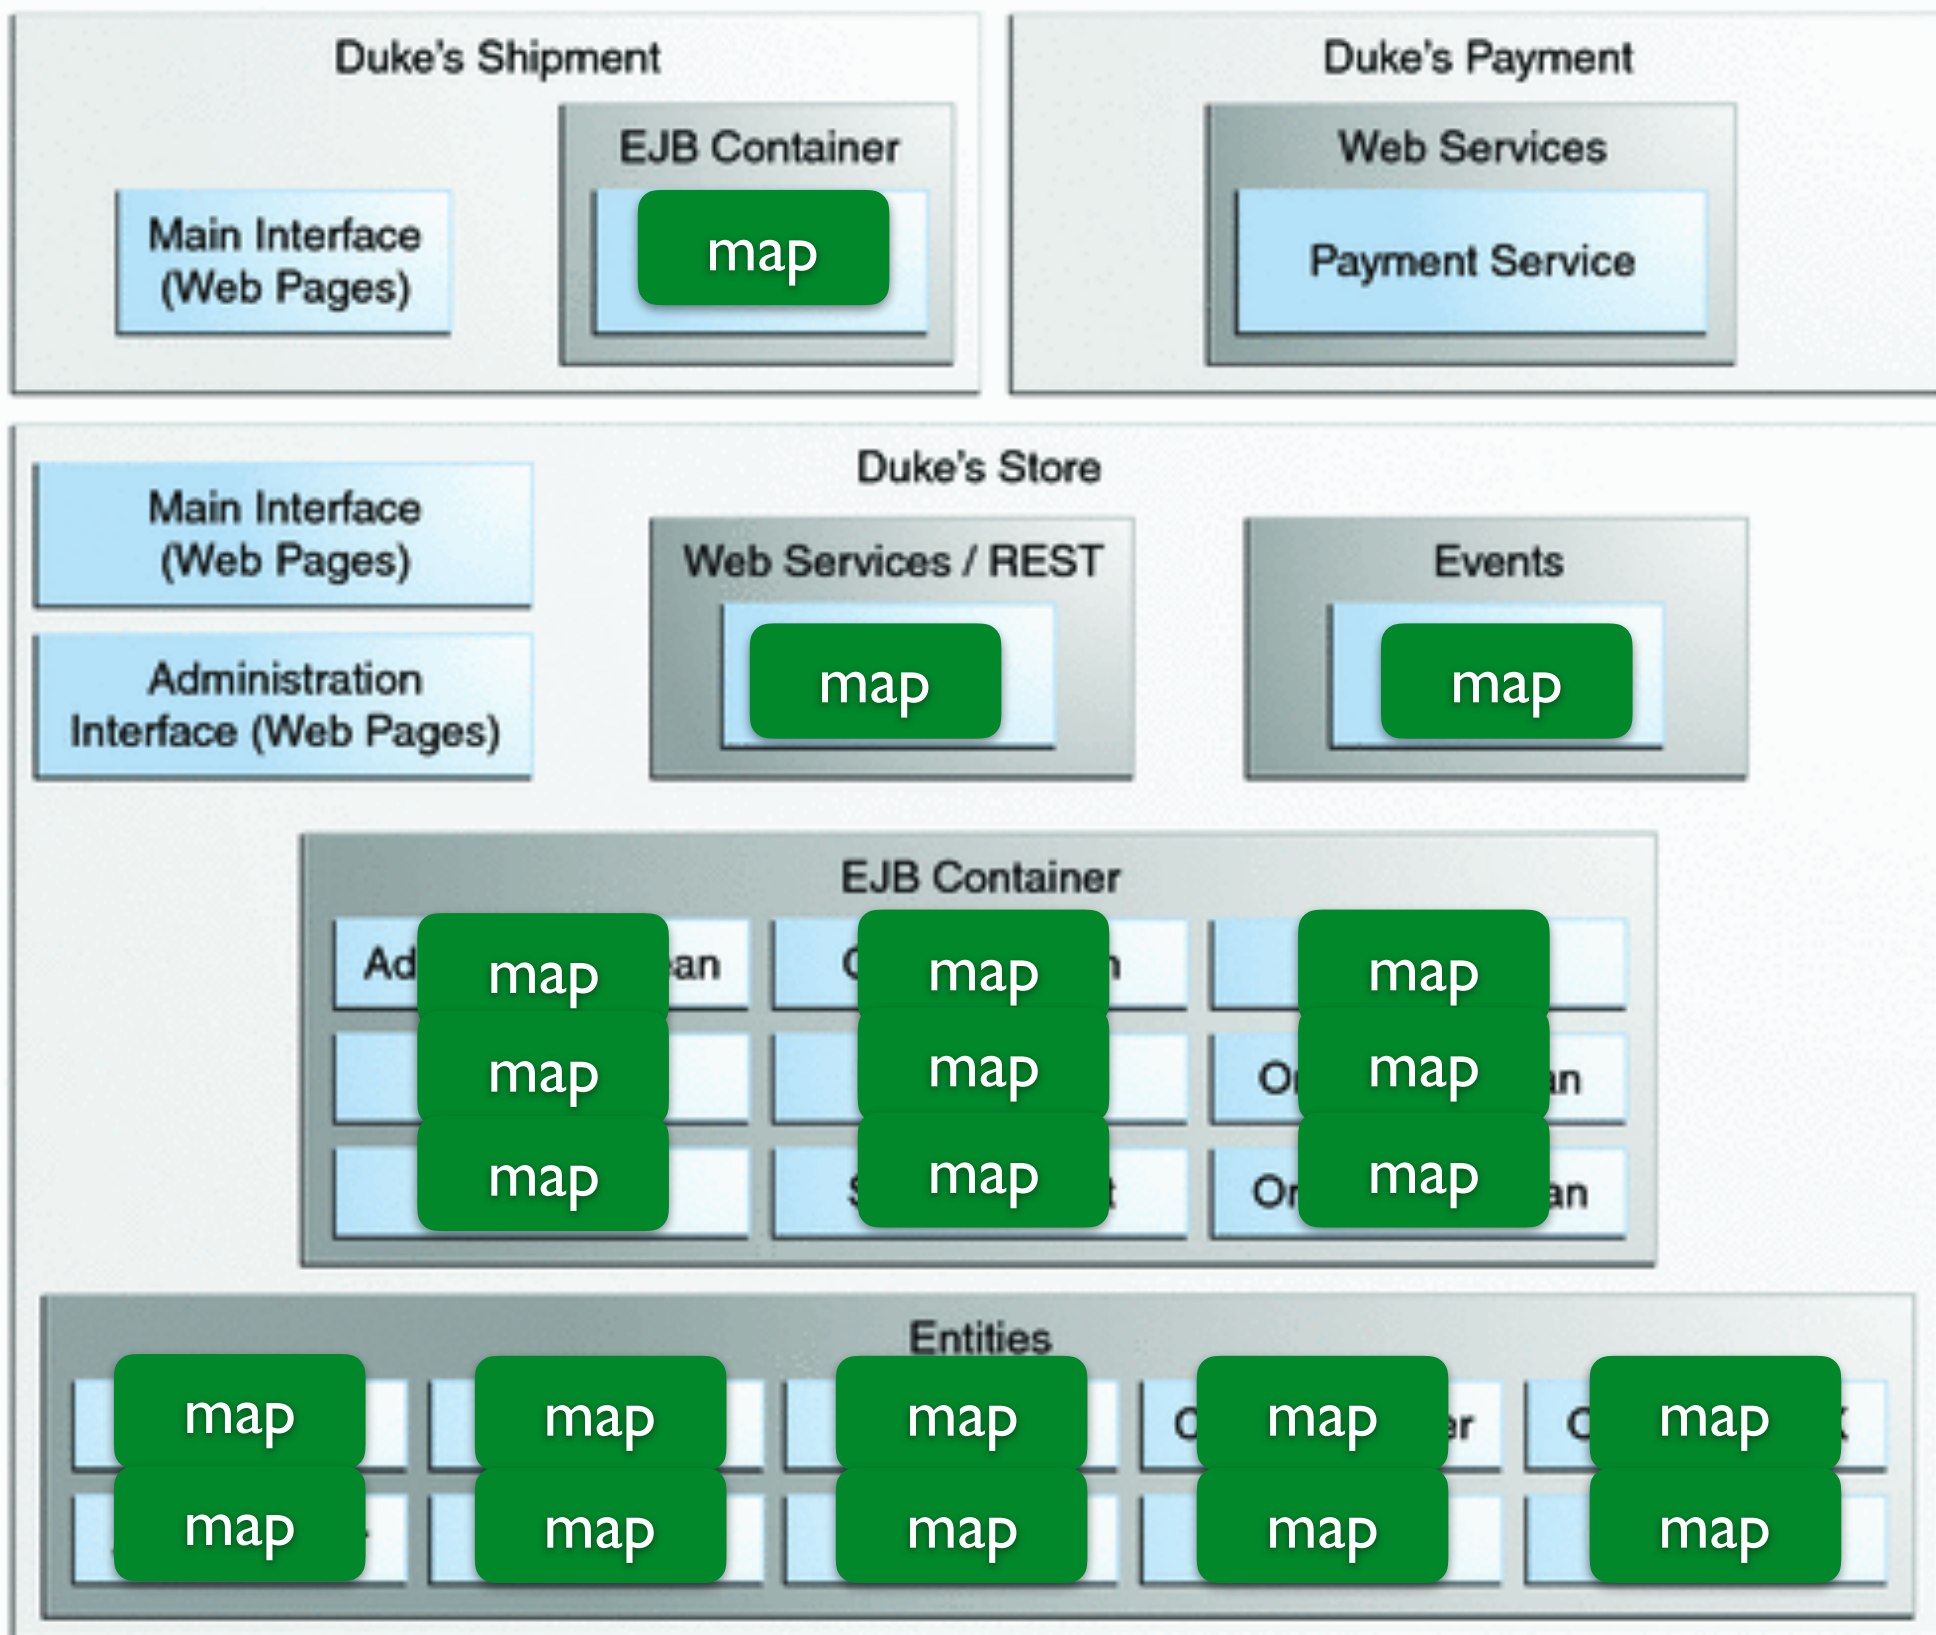
systemic

in libs, in apps, in config, on wires, at rest

specificity



generality



stability

semantic versioning is completely broken

we aspire to a different model

growth: accretion, relaxation, fixation

no breakage

The
Pragmatic
Programmers

Programming Clojure



2009

Stuart Halloway

Edited by Susannah Davidson Pfalzer

evidence for stability

Clojure release	date	breakage
1.0	05/2009	-
1.1	12/2009	-
1.2	08/2010	-
1.3	09/2011	1 LOC in test suite
1.4	04/2012	-
1.5	03/2013	-
1.6	03/2014	-
1.7	06/2015	-
1.8	01/2016	-
1.9	TBD	-

takeaways

build for yourself

value individuals *and* tools

value planning *and* agility

listen

trust yourself

will these ideas make
me a better
programmer?



Datomic

@stuarthalloway