



Datomic

A Day of Datomic



@stuarthalloway

Copyright Cognitect, Inc. All Rights Reserved.

agenda

what is Datomic?

information model

transaction model

query model

time model

operational model

what is Datomic?

a complete rethink of databases

agile

robust

powerful

time-aware

cloud

agile: universal schema

entity / attribute / value / tx / op

granular, attribute-level schema definition

easy to model

easy to migrate

Data Set: [e a v t added], 257 items				
e	a	v	t	added
:country/AF	:country/name	Afghanistan	13194139534626	true
:country/AL	:country/name	Albania	13194139534366	true
:country/DZ	:country/name	Algeria	13194139534392	true
:country/AS	:country/name	American Samoa	13194139534366	true
:country/AD	:country/name	Andorra	13194139534626	true

update in place

sharing	difficult
distribution	difficult
concurrent access	difficult
access pattern	eager
caching	difficult
examples	Java and .NET collections relational databases NoSQL databases

**NOT
ROBUST**

persistent data structures

sharing	trivial
distribution	easy
concurrent access	trivial
access pattern	eager or lazy
caching	easy
examples	Clojure, F# collections Datomic database

ROBUST

powerful

universal schema supports many access styles

row

column, key-value, document, and graph

declarative


datalog query

pull

time aware

db view	semantics	supports
<i>(default)</i>	current state	what is the current situation?
.asOf	state at point in past	how were things in the past?
.since	state since point in past	how have things changed?
tx report	before / after / change view of tx	automated event response
.with	state with proposed additions	what would happen if we did X?
.history	timeless view of all history	anything!

a database of little services



lab: running the tools

install Datomic, start dev transactor	http://docs.datomic.com/getting-started.html
install mbrainz dataset to datomic:dev://localhost:4334/mbrainz	https://github.com/Datomic/mbrainz-sample
explore mbrainz from console	http://docs.datomic.com/console.html
explore mbrainz from REST service	http://docs.datomic.com/rest.html
install examples	https://github.com/datomic/day-of-datomic
install your favorite REPL (links show Cursive)	https://www.jetbrains.com/idea/download/ + https://cursiveclojure.com/userguide/
play through CRUD example	https://github.com/Datomic/day-of-datomic/blob/master/tutorial/crud.clj

information model

notation

datoms

databases

entities

schema

notation: edn

```
{ :firstName "John"  
  :lastName "Smith"  
  :age 25  
  :address {  
    :streetAddress "21 2nd Street"  
    :city "New York"  
    :state "NY"  
    :postalCode "10021" }  
  :phoneNumber  
    [ { :type "name" :number "212 555-1234"}  
      { :type "fax" :number "646 555-4567" } ] }
```

type	examples
string	"foo"
character	\f
integer	42, 42N
floating point	3.14, 3.14M
boolean	true
nil	nil
symbol	foo, +
keyword	:foo, ::foo

type	properties	examples
list	sequential	(1 2 3)
vector	sequential and random access	[1 2 3]
map	associative	{ :a 100 :b 90 }
set	membership	# { :a :b }

generic extensibility

#name edn-form

name describes interpretation of following element

recursively defined

all data can be literal

built-in tags

#inst "rfc-3339-format"

tagged element is a string in RFC-3339 format

#uuid "f81d4fae-7dec-11d0-a765-00a0c91e6bf6"

tagged element is a canonical UUID string

datoms

granular, atomic facts

immutable

5-tuple: entity / attribute / value / transaction / op

example datoms

e	a	v	tx	op
jane	likes	broccoli	1008	true
jane	likes	pizza	1008	true
jane	likes	pizza	1148	false

datum API

```
public interface Datum {  
    Object e();  
    Object a();  
    Object v();  
    Object tx();  
    boolean added();  
  
    // positional  
    Object get(int index);  
}
```

database

a set of datoms (universal relation)

efficient storage

many sort orders

accumulate-only (not append-only)

database sorts in console

Query Entities Transactions **Indexes**

Index

A

V

E

T

Data Set: [e a v t added], 257 items

e	a	v	t	added
:country/AF	:country/name	Afghanistan	13194139534626	true
:country/AL	:country/name	Albania	13194139534366	true
:country/DZ	:country/name	Algeria	13194139534392	true
:country/AS	:country/name	American Samoa	13194139534366	true
:country/AD	:country/name	Andorra	13194139534626	true

entities

immutable

lazy

associative

inferred from datoms sharing a common **e**

point-in-time

bidirectional navigation

example entity

entity

```
{:db/id jane  
 :likes "broccoli"}
```



datoms

e	a	v	tx	op
jane	likes	broccoli	1008	true
jane	likes	pizza	1008	true
jane	likes	pizza	1148	false

entity API

```
interface Database
{
    Entity entity(Object entityId);
    // unrelated methods elided for brevity
}

interface Entity {
    Object get(Object key);
    Entity touch();
    Set keySet();
}
```

relation direction

```
// things that Jane likes  
jane.get("likes");
```

reversing direction

```
// people that like broccoli  
broccoli.get("_likes");
```



underscore means
“nav backwards”

entities in console

The screenshot shows a console interface with the 'Entities' tab selected. The 'Entity ID' field contains the value ':country/AL'. A red arrow points from the 'e' column of the 'Data Set' table to this field. The 'Data Set' table has columns 'e', 'a', 'v', 't', and 'added', with two rows: ':country/AF' (Afghanistan) and ':country/AL' (Albania). The ':country/AL' row is highlighted in blue.

Entity ID	Attribute	Value
:country/AL	:db/id	17592186045483
	:country/name	Albania
	:db/ident	:country/AL

e	a	v	t	added
:country/AF	:country/name	Afghanistan	13194139534626	true
:country/AL	:country/name	Albania	13194139534366	true

click any e

schema

schemas *add power*

schema is plain data

schema elements installed via transactions

make history-compatible changes at any time

common schema attributes

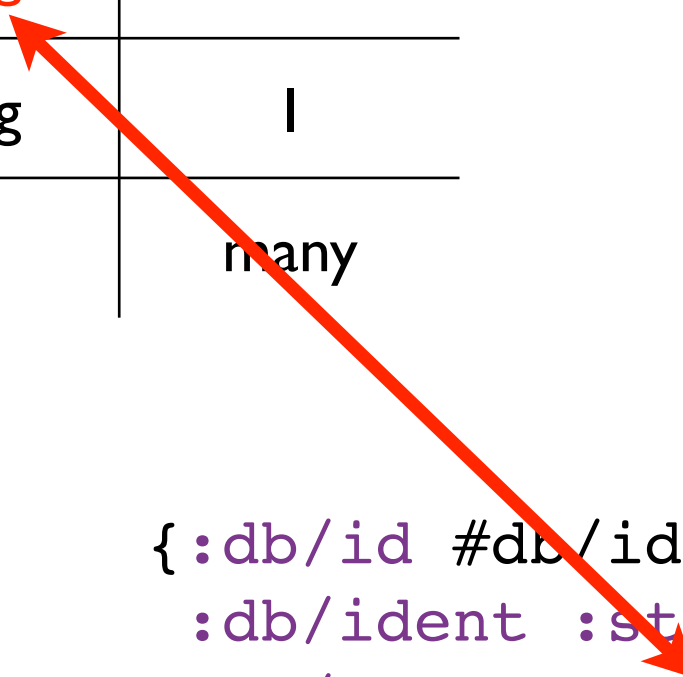
attribute	type	use
db/ident	keyword	programmatic name
db/valueType	ref	attribute type
db/cardinality	ref	one- or many- valued?
db/index	ref	creates AVET
db/unique	ref	unique, “upsert”
db/isComponent	ref	ownership

stories

attribute	type	cardinality
story/title	string	1
story/url	string	1
story/slug	string	1
news/comments	ref	many

schema is plain old data

attribute	type	card
story/title	string	1
story/url	string	1
story/slug	string	1
news/comments	ref	many



```
{:db/id #db/id[:db.part/db]  
 :db/ident :story/url  
 :db/valueType :db.type/string  
 :db/cardinality :db.cardinality/one  
 :db.install/_attribute :db.part/db}
```


users

attribute	type	cardinality
user/firstName	string	1
user/lastName	string	1
user/email*	string	1
user/upVotes	ref	many

*unique


comments

attribute	type	cardinality
comment/body	string	1
comment/author	ref	1
news/comments	ref	many

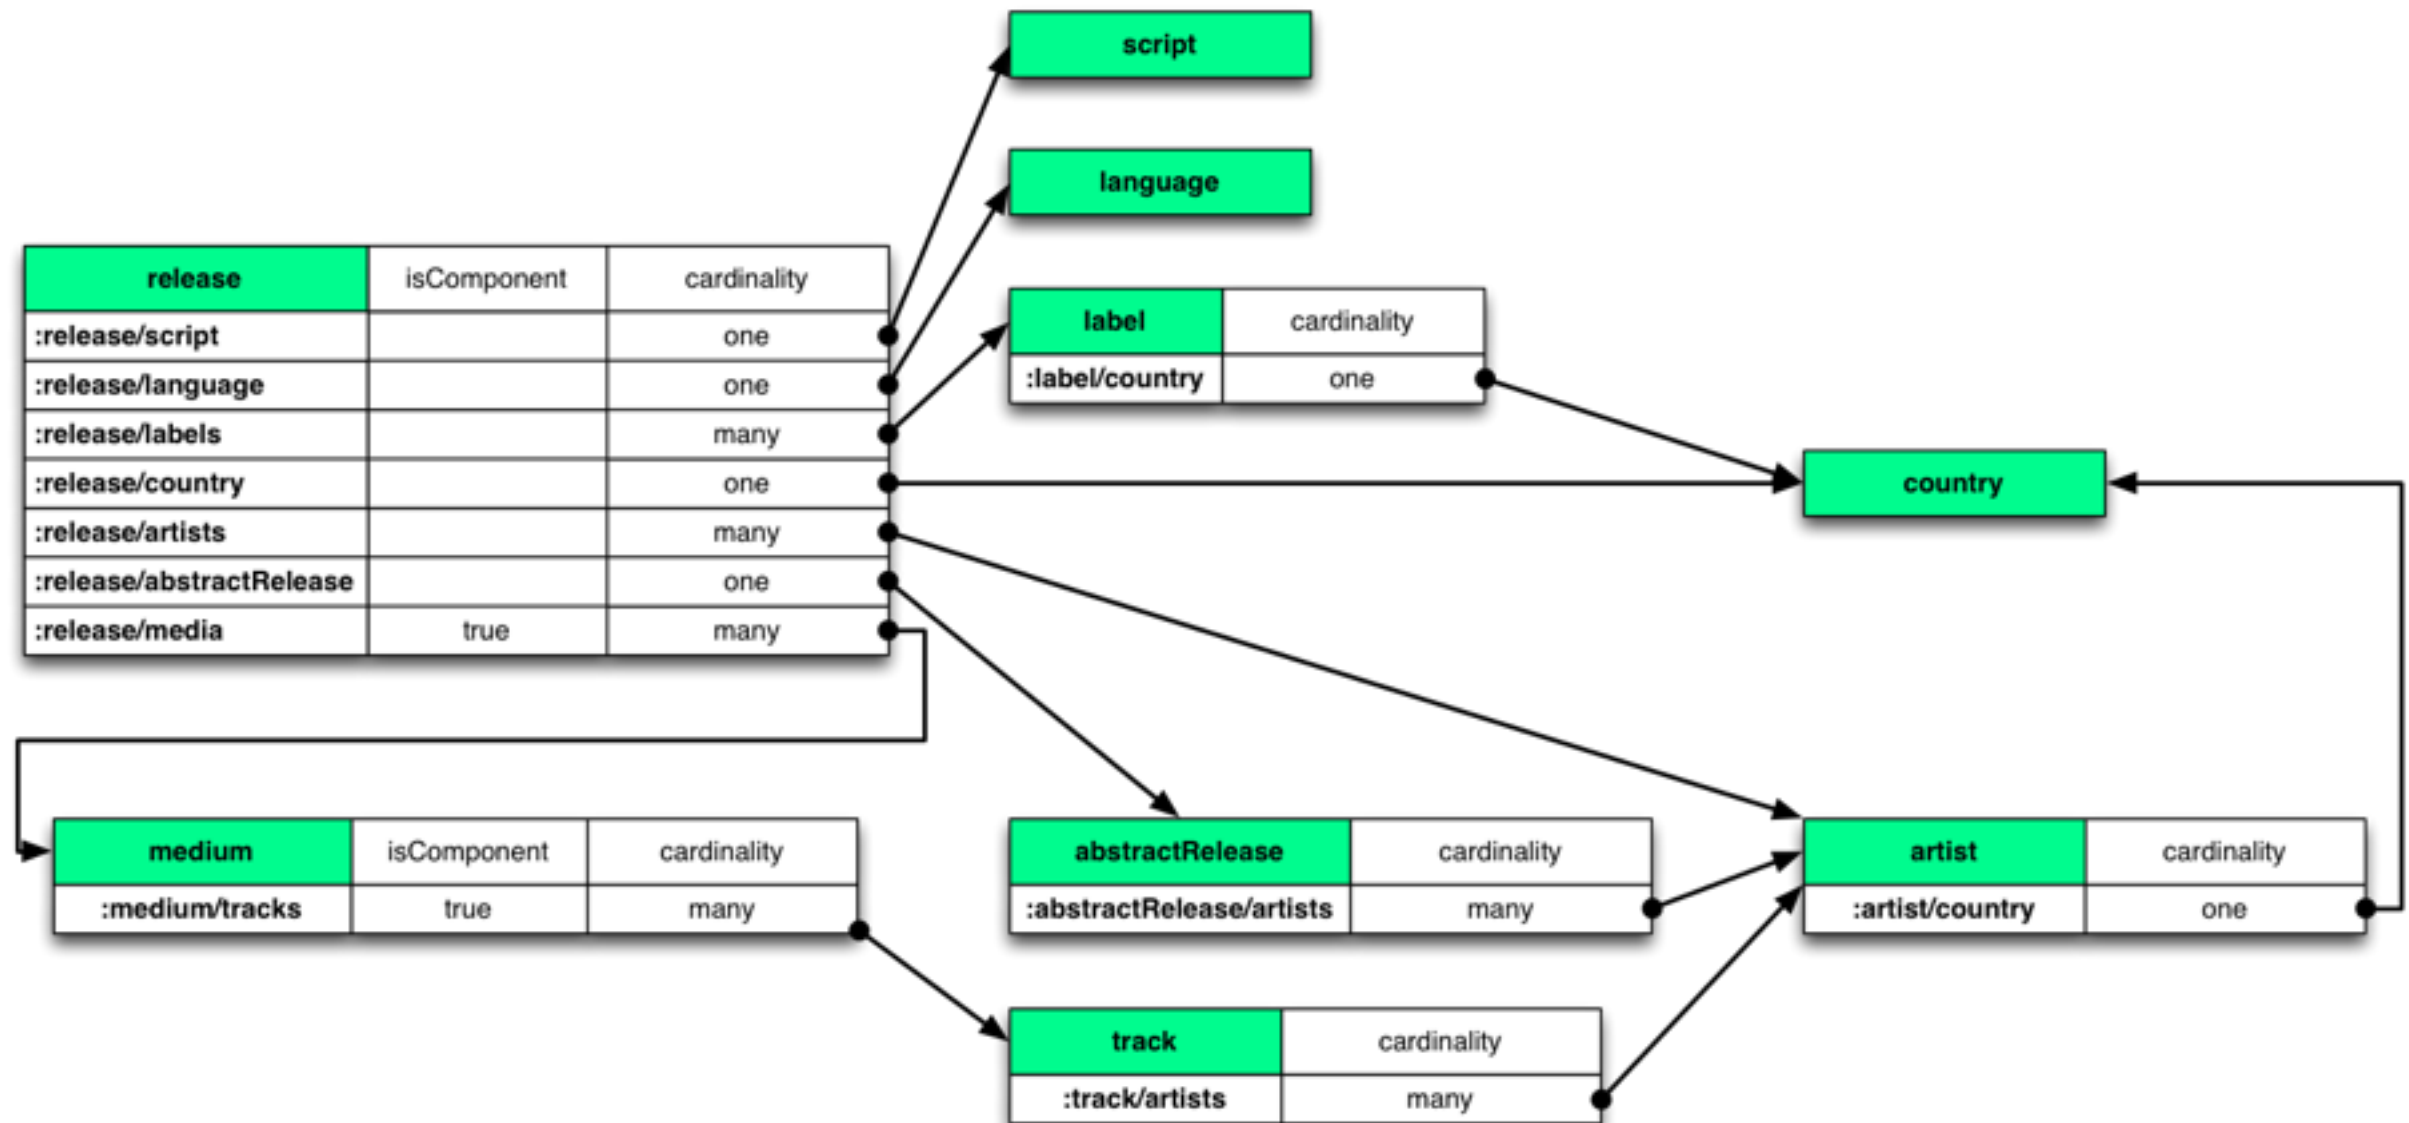
“types” do not dictate attrs

attribute	type	cardinality
story/title	string	1
story/url	string	1
story/slug	string	1
news/comments	ref	many

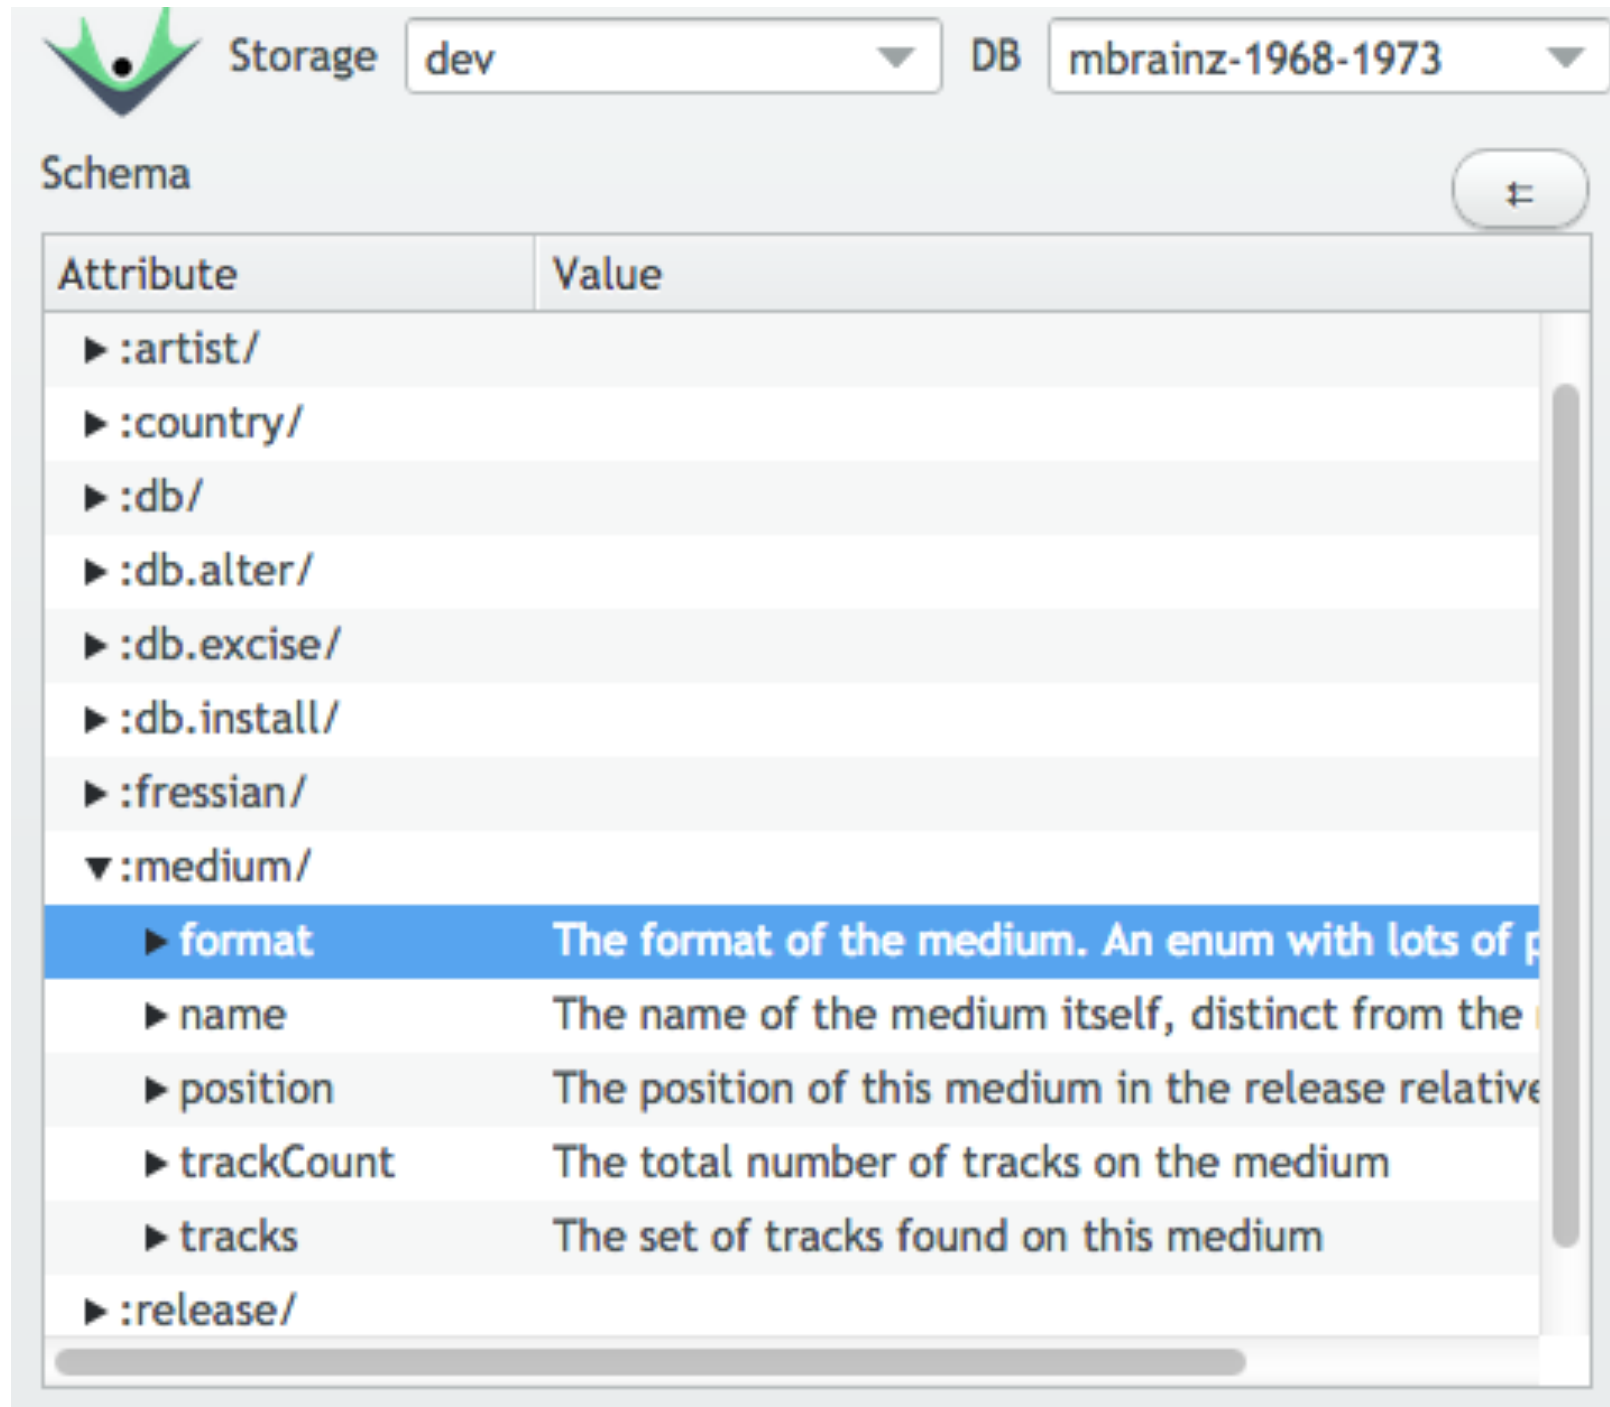
attribute	type	cardinality
comment/body	string	1
comment/author	ref	1
news/comments	ref	many



example: mbrainz



schema in console



Storage: dev DB: mbrainz-1968-1973

Schema

Attribute	Value
▶ :artist/	
▶ :country/	
▶ :db/	
▶ :db.alter/	
▶ :db.excise/	
▶ :db.install/	
▶ :fressian/	
▼ :medium/	
▶ format	The format of the medium. An enum with lots of p
▶ name	The name of the medium itself, distinct from the
▶ position	The position of this medium in the release relative
▶ trackCount	The total number of tracks on the medium
▶ tracks	The set of tracks found on this medium
▶ :release/	

lab: designing a schema

pick a **small** domain you know well

draw an entity/relationship diagram

convert to edn data

transact into a database

view schema in console

transaction model

ACID

assertion and retraction

entity maps

ids and partitions

uniqueness

transaction functions

ACID

atomic	transaction is a set of datoms transaction entirely in single write
consistent	all processes see same global ordering of transactions
isolated	single writer system (nobody to be isolated <i>from</i>)
durable	always flush through to durable storage before reporting transaction complete

assertion and retraction

```
[ :db/add entity-id attribute value]
```

```
[ :db/retract entity-id attribute value]
```

entity maps

concise form for multiple assertions about an entity


equivalent to corresponding list of asserts

can nest arbitrarily

lists vs. entity maps

```
[ :db/add 42 :likes "pizza" ]  
[ :db/add 42 :firstName "John" ]  
[ :db/add 42 :lastName "Doe" ]
```

assertion
list



```
{ :db/id 42  
  :likes "pizza"  
  :firstName "John"  
  :lastName "Doe" }
```

entity
map

cross reference

```
[ { :db/id #db/id[ :db.part/user -1 ]  
  :person/name "Bob"  
  :person/spouse #db/id[ :db.part/user -2 ] }
```

```
{ :db/id #db/id[ :db.part/user -2 ]  
  :person/name "Alice"  
  :person/spouse #db/id[ :db.part/user -1 ] } ]
```

nesting

```
[ { :db/id order-id  
  :order/lineItems [ { :lineItem/product chocolate  
                        :lineItem/quantity 1}  
                      { :lineItem/product whisky  
                        :lineItem/quantity 2 } ] } ] ]
```

partitions

partitions

group related entities in a partition

coarser granularity than e.g. tables

partition is a hint to indexing

group these things together

can help locality

does not affect semantics

partition locality

part	e	a	v	tx	added
0	0x0000000000000000	10	:db.part/db	0xc0000000000	TRUE
0	0x0000000000000000	11	0	0xc0000000036	TRUE
0	0x0000000000000000	11	3	0xc0000000036	TRUE
0	0x0000000000000000	11	4	0xc0000000036	TRUE
3	0x00000c0000000036	50	Wed Dec 31 19:00:00 E...	0xc0000000036	TRUE
3	0x00000c0000000038	50	Wed Dec 31 19:00:00 E...	0xc0000000038	TRUE
3	0x00000c000000003f	50	Wed Dec 31 19:00:00 E...	0xc000000003f	TRUE
3	0x00000c000000003e8	50	Mon Oct 13 18:52:59 E...	0xc00000003e8	TRUE
4	0x000010000000003eb	10	:country/BF	0xc00000003ea	TRUE
4	0x000010000000003eb	84	Burkina Faso	0xc00000003ea	TRUE
4	0x000010000000003ec	10	:country/JE	0xc00000003ea	TRUE
4	0x000010000000003ec	84	Jersey	0xc00000003ea	TRUE

built-in partitions

partition	usage
:db.part/db	schema entities
:db.part/tx	transaction entities
:db.part/user	user entities

creating partitions

```
[ { :db.install/_partition :db.part/db,  
  :db/id #db/id[:db.part/db],  
  :db/ident :inventory}  
{ :db.install/_partition :db.part/db,  
  :db/id #db/id[:db.part/db],  
  :db/ident :customers} ]
```

identity

identity

requirement	model with	value types
db-relative opaque id	entity id	opaque (long)
external id	:db.unique/identity attribute	string, uuid, uri
global opaque id	:db.unique/identity squuid	uuid
programmatic name	:db/ident	keyword

squids

semi-sequential UUIDs

do not fragment indexes

```
public class Peer;  
    public static UUID squuid();  
    public static long squuidTimeMillis(UUID squuid);  
    // other methods elided for brevity  
}
```

transaction functions

add and retract

```
[[:db/add john :likes pizza]  
[:db/retract john :likes iceCream]]
```

what about update?

```
[[:db/add john :likes pizza]  
[:db/retract john :likes iceCream]  
[:db/add john :balance 110?]]
```


atomic increment

```
[[:db/add john :likes pizza]  
[:db/retract john :likes iceCream]  
[:inc john :account 10]]
```

transaction fns

subset of data fns

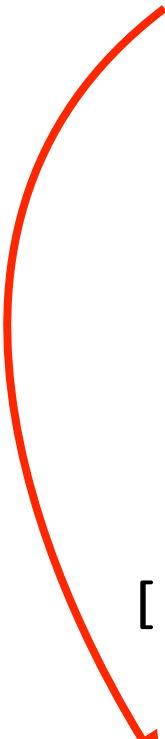
run inside transactions

have access to in-tx value of database

as first argument


tx function expansion

```
[[:db/add john :likes pizza]  
[:db/retract john :likes iceCream]  
[:inc john :balance 10]]
```



```
[[:db/add john :likes pizza]  
[:db/retract john :likes iceCream]  
[:db/add john :balance 110]]
```


lookup the function



```
[ :inc john :balance 10 ]
```

```
inc = db.entity( "inc" ).get( "db/fn" );
```


pass in current db




```
[ :inc john :account 10 ]
```

```
inc.invoke(db, ...);
```

pass in args




```
[ :inc john :account 10 ]
```



```
inc.invoke(db, john, :account, 10)
```

data out



```
[ :inc john :account 10 ]
```

```
inc.invoke(db, john, :account, 10)
```



```
[ [ :db/add john :account 110 ] ]
```

inc

```
public static Object inc(Object db, Object e, Object amount)
{
    // lookup entity
    // calculate new balance
    // create assertion
    // return list containing assertion
}
```


inc

```
public static Object inc(Object db, Object e, Object a, Object amount) {  
    Entity ent = ((Database)db).entity(e);  
    Long balance = (Long) ent.get(a) + (Long) amount;  
    List updated = list("db/add", e, a, balance);  
    return list(updated);  
}
```

lab: adding some data

create assertions and retractions	https://github.com/Datomic/day-of-datomic/blob/master/tutorial/crud.clj
create entity maps	https://github.com/Datomic/day-of-datomic/blob/master/tutorial/component_attributes.clj
modify an existing entity	https://github.com/Datomic/day-of-datomic/blob/master/tutorial/crud.clj
review your data at the console	
bonus: create a constructor function	https://github.com/Datomic/day-of-datomic/blob/master/tutorial/data_functions.clj

query model

datalog

pull

entities

raw indexes

why datalog?

equivalent to relational model + recursion

better fit than prolog for query

- no clause order dependency

- guaranteed termination

pattern-matching style easy to learn

example database

entity	attribute	value
42	:email	<u>jdoo@exampl.com</u>
43	:email	<u>jane@exampl.com</u>
42	:orders	107
42	:orders	141

data pattern

*Constrains the results returned,
binds variables*

```
[?customer :email ?email]
```

data pattern

*Constrains the results returned,
binds variables*

[?customer :email ?email]



entity



attribute



value

data pattern

*Constrains the results returned,
binds variables*

constant



[?customer :email ?email]

data pattern

*Constrains the results returned,
binds variables*

variable



variable



[?customer :email ?email]

example database

entity	attribute	value
42	:email	<u>jdoe@example.com</u>
43	:email	<u>jane@example.com</u>
42	:orders	107
42	:orders	141

[?customer :email ?email]

constants anywhere

“Find a particular customer’s email”

```
[ 42 :email ?email]
```

entity	attribute	value
42	:email	<u>jdoe@example.com</u>
43	:email	<u>jane@example.com</u>
42	:orders	107
42	:orders	141

[42 :email ?email]

variables anywhere

“What attributes does
customer 42 have?”

[42 **?attribute**]

entity	attribute	value
42	:email	<u>jdoe@example.com</u>
43	:email	<u>jane@example.com</u>
42	:orders	107
42	:orders	141

[42 ?attribute]

variables anywhere

“What attributes and values does
customer 42 have?”

[42 ?attribute ?value]

entity	attribute	value
42	:email	<u>jdoe@example.com</u>
43	:email	<u>jane@example.com</u>
42	:orders	107
42	:orders	141

[42 ?attribute ?value]

where clause


[:find ?customer
:where **[?customer :email]**]

data
pattern



find clause

variable to
return



```
[ :find ?customer  
  :where [?customer :email]]
```

implicit join

“Find all the customers who
have placed orders.”

```
[ :find ?customer  
  :where [ ?customer :email]  
          [ ?customer :orders ] ]
```

API

```
import static datomic.Peer.q;
```

```
q("[:find ?customer  
  :where [?customer :id]  
         [?customer :orders]]",  
  db);
```

q

```
import static datomic.Peer.q;
```

```
q( "[ :find ?customer  
      :where [?customer :id]  
              [?customer :orders] ] ",  
    db );
```

query

```
import static datomic.Peer.q;
```

```
q("[:find ?customer  
  :where [?customer :id]  
         [?customer :orders]]",  
  db);
```

input(s)

```
import static datomic.Peer.q;
```

```
q("[:find ?customer  
  :where [?customer :id]  
         [?customer :orders]]",  
  db);
```

in clause

*Names inputs so you can refer to them
elsewhere in the query*

```
:in $database ?email
```


parameterized query

“Find a customer by email.”

```
q([:find ?customer  
  :in $database ?email  
  :where [$database ?customer :email ?email]],  
db,  
"jdoe@example.com");
```

first input

“Find a customer by email.”

```
q([:find ?customer  
  :in $database ?email  
  :where [$database ?customer :email ?email]],  
  db,  
  "jdoe@example.com");
```

second input

“Find a customer by email.”

```
q([:find ?customer  
  :in $database ?email  
  :where [$database ?customer :email ?email]],  
db,  
"jdoe@example.com");
```

verbose?

“Find a customer by email.”

```
q([:find ?customer  
  :in $database ?email  
  :where [$database ?customer :email ?email]],  
  db,  
  "jdoe@example.com");
```

shortest name possible

“Find a customer by email.”

```
q([:find ?customer  
  :in $ ?email  
  :where [$ ?customer :email ?email]],  
  db,  
  "jdoe@example.com");
```

elide \$ in where

“Find a customer by email.”

```
q([:find ?customer  
  :in $ ?email  
  :where [ ?customer :email ?email]],  
db,  
"jdoe@example.com");
```



no need to
specify \$

predicates

*Functional constraints that can
appear in a :where clause*

```
[ (< 50 ?price) ]
```

adding a predicate

“Find the expensive items”

```
[ :find ?item  
  :where [?item :item/price ?price]  
          [ (< 50 ?price) ] ]
```


functions

*Take bound variables as inputs
and bind variables with output*

```
[ (shipping ?zip ?weight) ?cost]
```

function args


[(shipping ?zip ?weight) ?cost]



bound inputs

function returns

```
[ (shipping ?zip ?weight) ?cost ]
```



bind return
values

calling a function

“Find me the customer/product combinations where the shipping cost dominates the product cost.”

```
[ :find ?customer ?product
  :where [?customer :shipAddress ?addr]
          [?addr :zip ?zip]
          [?product :product/weight ?weight]
          [?product :product/price ?price]
          [(Shipping/estimate ?zip ?weight) ?shipCost]
          [(<= ?price ?shipCost)]]
```

← navigate from customer to zip

calling a function

“Find me the customer/product combinations where the shipping cost dominates the product cost.”

```
[ :find ?customer ?product
  :where [?customer :shipAddress ?addr]
          [?addr :zip ?zip]
          [?product :product/weight ?weight]
          [?product :product/price ?price]
          [(Shipping/estimate ?zip ?weight) ?shipCost]
          [(<= ?price ?shipCost)]]
```

get product facts
needed *during query*

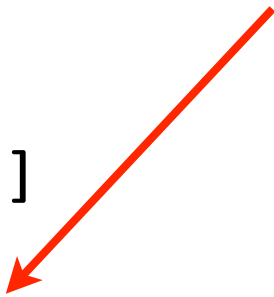


calling a function

“Find me the customer/product combinations where the shipping cost dominates the product cost.”

```
[ :find ?customer ?product
  :where [?customer :shipAddress ?addr]
         [?addr :zip ?zip]
         [?product :product/weight ?weight]
         [?product :product/price ?price]
         [(Shipping/estimate ?zip ?weight) ?shipCost]
         [(<= ?price ?shipCost)]]
```

call web service
to bind shipCost



BYO functions

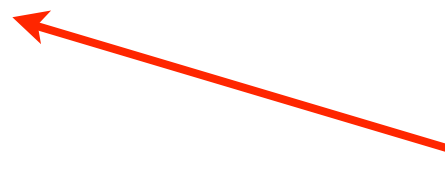
*Functions can be plain
JVM code.*

```
public class Shipping {  
    public static BigDecimal  
    estimate(String zip1, int pounds);  
}
```

calling a function

“Find me the customer/product combinations where the shipping cost dominates the product cost.”

```
[ :find ?customer ?product
  :where [?customer :shipAddress ?addr]
          [?addr :zip ?zip]
          [?product :product/weight ?weight]
          [?product :product/price ?price]
          [(Shipping/estimate ?zip ?weight) ?shipCost]
          [(<= ?price ?shipCost)]]
```



constrain price

calling a function

“Find me the customer/product combinations where the shipping cost dominates the product cost.”

```
[ :find ?customer ?product  
  :where [?customer :shipAddress ?addr]  
          [?addr :zip ?zip]  
          [?product :product/weight ?weight]  
          [?product :product/price ?price]  
          [(Shipping/estimate ?zip ?weight) ?shipCost]  
          [(<= ?price ?shipCost)]]
```

← return customer,
product pairs

calling a function

“Find me the customer/product combinations where the shipping cost dominates the product cost.”


```
[ :find ?customer ?product
  :where [?customer :shipAddress ?addr]
          [?addr :zip ?zip]
          [?product :product/weight ?weight]
          [?product :product/price ?price]
          [(Shipping/estimate ?zip ?weight) ?shipCost]
          [(<= ?price ?shipCost)]]
```

entities

maplike, point-in-time view of datoms sharing a common **e**

```
{:db/id 42  
 :likes "pizza"  
 :firstName "John"  
 :lastName "Doe"}
```

entity



datoms

```
[42 :likes "pizza"]  
[42 :firstName "John"]  
[42 :lastName "Doe"]
```

entities

transformation is purely mechanical

`{ :db/id 42` ← special key for e
 `:likes "pizza"`
 `:firstName "John"`
 `:lastName "Doe" }`

```
[ 42 :likes "pizza" ]  
[ 42 :firstName "John" ]  
[ 42 :lastName "Doe" ]
```

one database, many indexes

structure	attribute
k/v	AVET
row	EAVT
column	AEVT
document	EAVT, partitions, components
graph	VAET

lab: query

explore mbrainz data with Query API	https://github.com/Datomic/mbrainz-sample/wiki/Queries
explore mbrainz data with Pull API	http://docs.datomic.com/pull.html
query your own data	http://docs.datomic.com/query.html
pull your own data	http://docs.datomic.com/pull.html
navigate your data with the Entity API	http://docs.datomic.com/entities.html

time model

reified transactions

t & basis-t

log

filters

sync

excision

reified transactions

entities like any other entity in the system

associated with every datom

increasing entity ids over time

associated with increasing counter **t**

have a `:db/txInstant`

have any other attributes you specify

have their own index (the log)

basis

```
(def basis-t (d/basis-t db))  
=> 130223
```

```
(def basis-tx (d/t->tx basis-t))  
=> 13194139663535
```

```
(d/pull db '[*] basis-tx)  
=> {:db/id 13194139663535,  
    :db/txInstant #inst "2014-10-13T22:58:03.832-00:00"}
```

basis is most recent t, tx



facts about basis tx

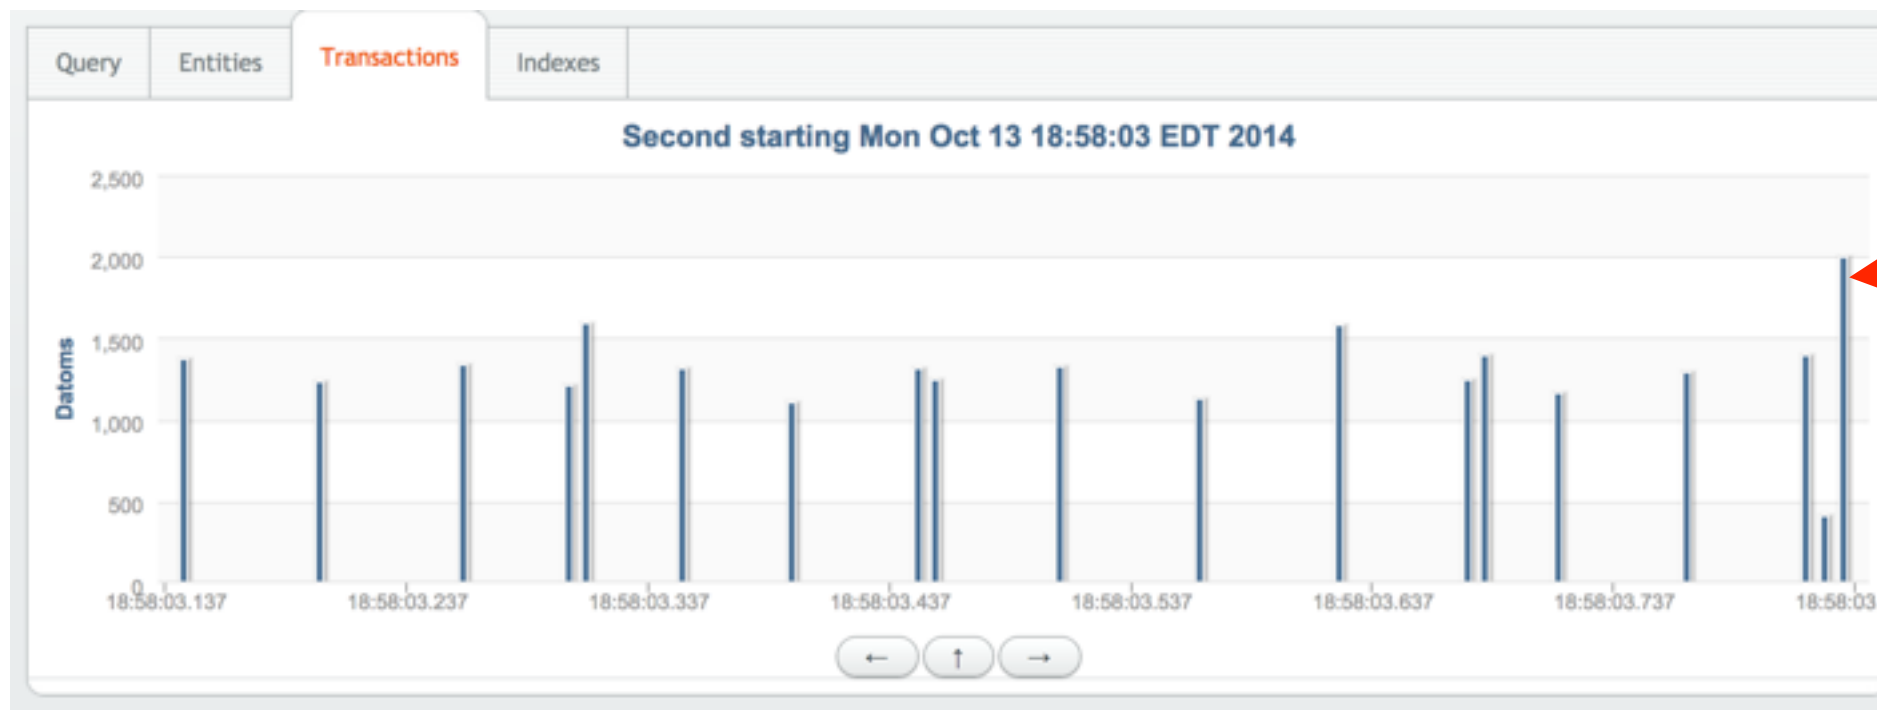


a time point is any of: t, tx, instant

log

```
(def log (d/log conn))  
(-> (d/tx-range log basis-tx (inc basis-tx))  
    seq first :data count)
```



contents of txes



transaction attributes

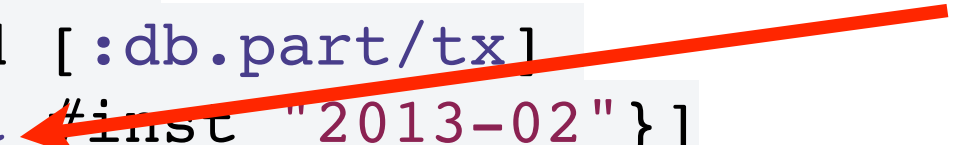
tx partition
references current tx

```
[{:db/id #db/id [:db.part/user]  
  :story/title "codeq"  
  :story/url "http://blog.datomic.com/2012/10/codeq.html"}  
{:db/id (d/tempid :db.part/tx)  
  :publish/at (java.util.Date.)}]
```



add your own
attributes

```
[{:db/id [:item/id "DLC-042"]  
  :item/count 250}  
{:db/id #db/id [:db.part/tx]  
  :db/txInstant #inst "2013-02"}]
```



override txInstant
(for imports)

<https://github.com/Datomic/day-of-datomic/blob/master/tutorial/filter.clj>

<https://github.com/Datomic/day-of-datomic/blob/master/tutorial/filters.clj>

filters

filter	semantics	supports
<i>(default)</i>	current state	what is the current situation?
<code>.asOf</code>	state at point in past	how were things in the past?
<code>.since</code>	state since point in past	how have things changed?
<code>.history</code>	timeless view of all history	anything!

raw data

e	a	v	tx	added
0x0000c00000003e9	:db/txInstant	Mon Dec 31 19:00:00	0xc00000003e9	TRUE
0x00001000000003ea	:item/id	DLC-042	0xc00000003e9	TRUE
0x00001000000003ea	:item/description	Dilitihium Crystals	0xc00000003e9	TRUE
0x00001000000003ea	:item/count	100	0xc00000003e9	TRUE
0x0000c00000003eb	:db/txInstant	Thu Jan 31 19:00:00	0xc00000003eb	TRUE
0x00001000000003ea	:item/count	100	0xc00000003eb	FALSE
0x00001000000003ea	:item/count	250	0xc00000003eb	TRUE
0x0000c00000003ec	:db/txInstant	Thu Feb 27 19:00:00	0xc00000003ec	TRUE
0x00001000000003ea	:item/count	250	0xc00000003ec	FALSE
0x00001000000003ea	:item/count	50	0xc00000003ec	TRUE
0x0000c00000003ed	:db/txInstant	Mon Mar 31 20:00:00	0xc00000003ed	TRUE
0x0000c00000003ed	:tx/error	TRUE	0xc00000003ed	TRUE
0x00001000000003ea	:item/count	50	0xc00000003ed	FALSE
0x00001000000003ea	:item/count	9999	0xc00000003ed	TRUE
0x0000c00000003ee	:db/txInstant	Wed May 14 20:00:00	0xc00000003ee	TRUE
0x00001000000003ea	:item/count	9999	0xc00000003ee	FALSE
0x00001000000003ea	:item/count	100	0xc00000003ee	TRUE

default filter

e	a	v	tx	added
0x0000c00000003e9	:db/txInstant	Mon Dec 31 19:00:00	0xc00000003e9	TRUE
0x00001000000003ea	:item/id	DLC-042	0xc00000003e9	TRUE
0x00001000000003ea	:item/description	Dilitihium Crystals	0xc00000003e9	TRUE
0x00001000000003ea	:item/count	100	0xc00000003e9	TRUE
0x0000c00000003eb	:db/txInstant	Thu Jan 31 19:00:00	0xc00000003eb	TRUE
0x00001000000003ea	:item/count	100	0xc00000003eb	FALSE
0x00001000000003ea	:item/count	250	0xc00000003eb	TRUE
0x0000c00000003ec	:db/txInstant	Thu Feb 27 19:00:00	0xc00000003ec	TRUE
0x00001000000003ea	:item/count	250	0xc00000003ec	FALSE
0x00001000000003ea	:item/count	50	0xc00000003ec	TRUE
0x0000c00000003ed	:db/txInstant	Mon Mar 31 20:00:00	0xc00000003ed	TRUE
0x0000c00000003ed	:tx/error	TRUE	0xc00000003ed	TRUE
0x00001000000003ea	:item/count	50	0xc00000003ed	FALSE
0x00001000000003ea	:item/count	9999	0xc00000003ed	TRUE
0x0000c00000003ee	:db/txInstant	Wed May 14 20:00:00	0xc00000003ee	TRUE
0x00001000000003ea	:item/count	9999	0xc00000003ee	FALSE
0x00001000000003ea	:item/count	100	0xc00000003ee	TRUE

as-of filter

e	a	v	tx	added
0x0000c00000003e9	:db/txInstant	Mon Dec 31 19:00:00	0xc00000003e9	TRUE
0x00001000000003ea	:item/id	DLC-042	0xc00000003e9	TRUE
0x00001000000003ea	:item/description	Dilitihium Crystals	0xc00000003e9	TRUE
0x00001000000003ea	:item/count	100	0xc00000003e9	TRUE
0x0000c00000003eb	:db/txInstant	Thu Jan 31 19:00:00	0xc00000003eb	TRUE
0x00001000000003ea	:item/count	100	0xc00000003eb	FALSE
0x00001000000003ea	:item/count	250	0xc00000003eb	TRUE
0x0000c00000003ec	:db/txInstant	Thu Feb 27 19:00:00	0xc00000003ec	TRUE
0x00001000000003ea	:item/count	250	0xc00000003ec	FALSE
0x00001000000003ea	:item/count	50	0xc00000003ec	TRUE
0x0000c00000003ed	:db/txInstant	Mon Mar 31 20:00:00	0xc00000003ed	TRUE
0x0000c00000003ed	:tx/error	TRUE	0xc00000003ed	TRUE
0x00001000000003ea	:item/count	50	0xc00000003ed	FALSE
0x00001000000003ea	:item/count	9999	0xc00000003ed	TRUE
0x0000c00000003ee	:db/txInstant	Wed May 14 20:00:00	0xc00000003ee	TRUE
0x00001000000003ea	:item/count	9999	0xc00000003ee	FALSE
0x00001000000003ea	:item/count	100	0xc00000003ee	TRUE


since filter

e	a	v	tx	added
0x0000c00000003e9	:db/txInstant	Mon Dec 31 19:00:00	0xc00000003e9	TRUE
0x00001000000003ea	:item/id	DLC-042	0xc00000003e9	TRUE
0x00001000000003ea	:item/description	Dilitihium Crystals	0xc00000003e9	TRUE
0x00001000000003ea	:item/count	100	0xc00000003e9	TRUE
0x0000c00000003eb	:db/txInstant	Thu Jan 31 19:00:00	0xc00000003eb	TRUE
0x00001000000003ea	:item/count	100	0xc00000003eb	FALSE
0x00001000000003ea	:item/count	250	0xc00000003eb	TRUE
0x0000c00000003ec	:db/txInstant	Thu Feb 27 19:00:00	0xc00000003ec	TRUE
0x00001000000003ea	:item/count	250	0xc00000003ec	FALSE
0x00001000000003ea	:item/count	50	0xc00000003ec	TRUE
0x0000c00000003ed	:db/txInstant	Mon Mar 31 20:00:00	0xc00000003ed	TRUE
0x0000c00000003ed	:tx/error	TRUE	0xc00000003ed	TRUE
0x00001000000003ea	:item/count	50	0xc00000003ed	FALSE
0x00001000000003ea	:item/count	9999	0xc00000003ed	TRUE
0x0000c00000003ee	:db/txInstant	Wed May 14 20:00:00	0xc00000003ee	TRUE
0x00001000000003ea	:item/count	9999	0xc00000003ee	FALSE
0x00001000000003ea	:item/count	100	0xc00000003ee	TRUE

history = raw

e	a	v	tx	added
0x0000c00000003e9	:db/txInstant	Mon Dec 31 19:00:00	0xc00000003e9	TRUE
0x00001000000003ea	:item/id	DLC-042	0xc00000003e9	TRUE
0x00001000000003ea	:item/description	Dilitihium Crystals	0xc00000003e9	TRUE
0x00001000000003ea	:item/count	100	0xc00000003e9	TRUE
0x0000c00000003eb	:db/txInstant	Thu Jan 31 19:00:00	0xc00000003eb	TRUE
0x00001000000003ea	:item/count	100	0xc00000003eb	FALSE
0x00001000000003ea	:item/count	250	0xc00000003eb	TRUE
0x0000c00000003ec	:db/txInstant	Thu Feb 27 19:00:00	0xc00000003ec	TRUE
0x00001000000003ea	:item/count	250	0xc00000003ec	FALSE
0x00001000000003ea	:item/count	50	0xc00000003ec	TRUE
0x0000c00000003ed	:db/txInstant	Mon Mar 31 20:00:00	0xc00000003ed	TRUE
0x0000c00000003ed	:tx/error	TRUE	0xc00000003ed	TRUE
0x00001000000003ea	:item/count	50	0xc00000003ed	FALSE
0x00001000000003ea	:item/count	9999	0xc00000003ed	TRUE
0x0000c00000003ee	:db/txInstant	Wed May 14 20:00:00	0xc00000003ee	TRUE
0x00001000000003ea	:item/count	9999	0xc00000003ee	FALSE
0x00001000000003ea	:item/count	100	0xc00000003ee	TRUE

sync



excision: permanently,
unrecoverably, lose data

don't do this
unless you must

lab: navigating time

explore your data with filters in console	
explore your data with transaction view in console	
query against an as-of db	https://github.com/Datomic/day-of-datomic/blob/master/tutorial/filters.clj
use the log to recover a recent transaction	https://github.com/Datomic/day-of-datomic/blob/master/tutorial/basis_t_and_log.clj
annotate a transaction	https://github.com/Datomic/day-of-datomic/blob/master/tutorial/filters.clj

operational model

caching

indexing

bring your own storage


transactors

peers

monitoring

capacity planning

caching



understanding caching

single knob: how much?

segments (1000s of datoms), not individual facts

segments are immutable, never expire

segment names can cohabit with other caches


object cache is *in your process!*

will beat any RPC-based database when cache hits


reading data

source	format	capacity	latency
peer object cache	Java objects	~ 10 GB / peer	sub-microsecond
memcached	zipped fressian	arbitrary	~ 1 msec
storage	zipped fressian	arbitrary	~ 10 msec

indexing



index trees



segments fit in cache for small dbs,
1-2 reads away for larger dbs

*history is in separate trees, and is
never in the way of present queries*

understanding indexing

indexes made as needed


background process (on transactor)

indexing can fall too far behind

memory index exceeds predefined threshold

transactor throttles writes

only likely to see during import jobs



storage options

dynamodb

sql (any JDBC)

cassandra

riak

couchbase

infinispan

dev (local disk)

choosing storage

more important

reliability


manageability

familiarity

less important

datomic insulates storage from load

datomic tolerates storage latency



understanding transactors

ACID (single writer thread!)

stream recent change to peers

manage indexing


dead-simple HA configuration

point two machines at one storage

dead-simple HA coordination

conditional puts on storage

peers



understanding peers

embedded JVM lib

directly in application servers (most common)

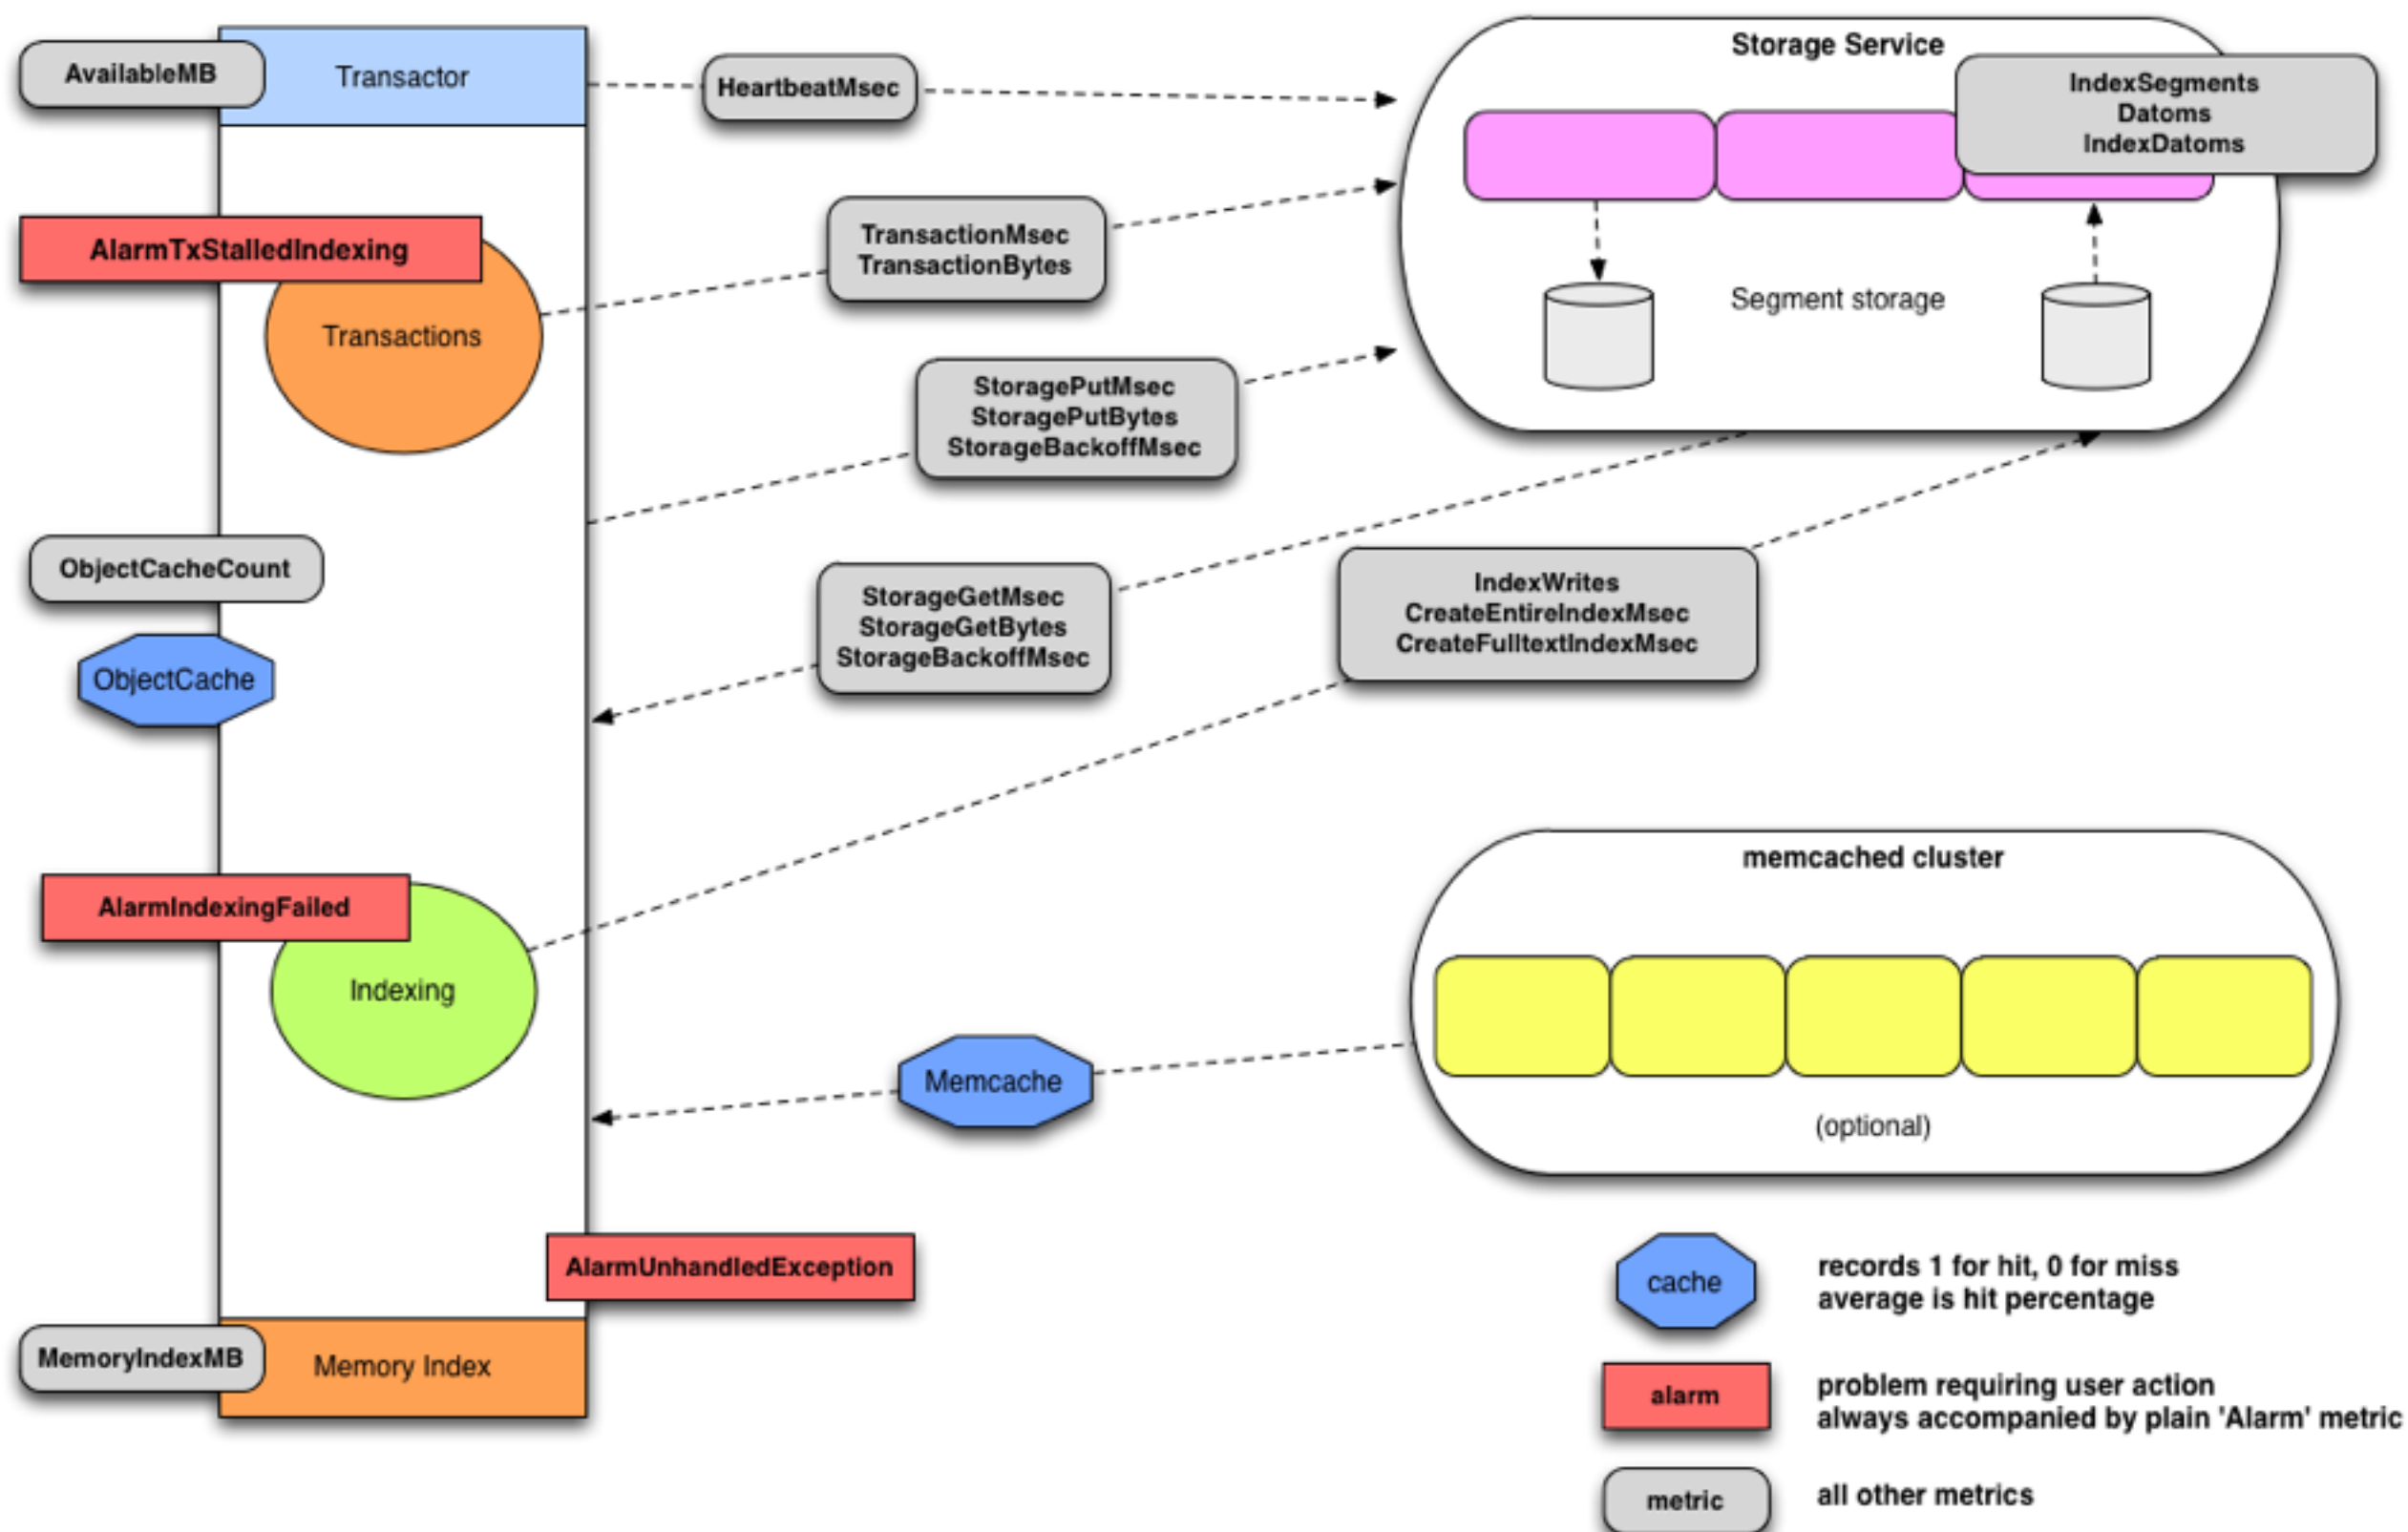
introduce a service tier

query load does not bother the transactor

answer queries in-memory via automatic cache

`datomic.objectCacheMax` (default 50% RAM)

Datomic Monitoring




LogWriteMsec (Milliseconds)

Average ▾

5 Minutes ▾



healthy transactor
log writes




Left axis units: Milliseconds

LogWriteMsec

AlarmBackPressure (Count)

Sum ▾

5 Minutes ▾



Left axis units: Count

AlarmBackPressure

capacity planning

data size	up to 10 billion datoms
write volume	consistent with data size in the long run
read volume	elastically scale peers as needed
housekeeping	http://docs.datomic.com/capacity.html#garbage-collection
caching	test and monitor

Datomic is a poor fit for

write scale

media storage (unstructured docs, audio, video)

churn (e.g. hit counter)

*big data, graphics, video, unstructured documents, log
file analysis*

Datomic is a good fit for

valuable information of record

dev and ops flexibility

history

read scale

transactional data, business records, medical records, financial records, scientific records, inventory, configuration, web apps, departmental databases, cloud apps

lab: assess Datomic

for your dataset, create a comparison matrix

which Datomic characteristics are

beneficial?

detrimental?

irrelevant?

compared to ?

10 billion datoms
access pattern at read-time
ACID
as-of
attributes
background indexing
cardinality
cloudwatch monitoring
column access (AEVT)
component attributes
cross-db joins
datalog
data API
document access
edn

elastically scalable read
entity maps
excision
graph access (VAET)
high availability
history
joins
lazy entities
lookup refs
memory speeds
partitions
pluggable monitoring
pluggable storage
pull
query functions
query predicates

raw index access
reified transactions
row access (EAVT)
rules
since
sync
time points
transaction functions
transparent memcached
tx reports
unique ids
universal schema
value access (AVET)
upsert
with

thanks!

@stuarthalloway

<https://github.com/stuarthalloway/presentations/wiki>

<http://www.linkedin.com/pub/stu-halloway/0/110/543/>

<https://twitter.com/stuarthalloway>

<mailto:stu@cognitect.com>