

making Datomic

@stuarthalloway

making things

Relevance

Clojure

Datomic

ClojureScript

core.async

Pedestal

Cognitect



Datomic

functional

well-known benefits

referential transparency, composition, testing

concurrency friendly

less obvious benefits

time model

topology

agility

flexibility

universal schema

programming with data

dynamic

power

indexes

logic

ACID

functional, lazy peers

```
Connection conn =  
connect("datomic:ddb://us-east-1/mb/mbrainz");
```

```
Database db = conn.db();
```

```
Set results = q(..., db);
```

```
Set crossDbResults = q(..., db1, db2);
```

```
Entity e = db.entity(42);
```

functional, lazy peers

```
Connection conn =  
connect( "datomic:ddb://us-east-1/mb/mbbrainz" );
```

← pluggable storage protocol

```
Database db = conn.db();
```

```
Set results = q(..., db);
```

```
Set crossDbResults = q(..., db1, db2);
```

```
Entity e = db.entity(42);
```

functional, lazy peers

```
Connection conn =  
connect("datomic:ddb://us-east-1/mb/mbbrainz");
```

```
Database db = conn.db();
```



database is a lazily
realized value, available
to all peers equally

```
Set results = q(..., db);
```

```
Set crossDbResults = q(..., db1, db2);
```

```
Entity e = db.entity(42);
```


functional, lazy peers

```
Connection conn =  
connect("datomic:ddb://us-east-1/mb/mbrainz");
```

```
Database db = conn.db();
```

```
Set results = q(..., db);
```



query databases,
not connections

```
Set crossDbResults = q(..., db1, db2);
```

```
Entity e = db.entity(42);
```

functional, lazy peers

```
Connection conn =  
connect( "datomic:ddb://us-east-1/mb/mbrainz" );
```

```
Database db = conn.db( );
```

```
Set results = q( ..., db );
```

```
Set crossDbResults = q( ..., db1, db2 );
```

```
Entity e = db.entity(42);
```



join across databases,
systems, in-memory collections

functional, lazy peers

```
Connection conn =  
connect("datomic:ddb://us-east-1/mb/mbrainz");
```

```
Database db = conn.db();
```

```
Set results = q(..., db);
```

```
Set crossDbResults = q(..., db1, db2);
```

```
Entity e = db.entity(42);
```




lazy, associative
navigable value

ACID, serialized, time aware

```
List newData = ...;  
Future<Map> f = conn.transactAsync(list);  
  
dbBefore = conn.db.asOf(time);  
  
possibleFuture = db.with(...);  
  
allTime = db.history();  
  
BlockingQueue<Map> queue = conn.txReportQueue();  
  
Log log = conn.log();  
Iterable<Map> it = log.txRange(startOfMonth, null);
```

ACID, serialized, time aware

```
List newData = ...;  
Future<Map> f = conn.transactAsync(list);  
  
dbBefore = conn.db.asOf(time);  
possibleFuture = db.with(...);  
  
allTime = db.history();  
  
BlockingQueue<Map> queue = conn.txReportQueue();  
  
Log log = conn.log();  
Iterable<Map> it = log.txRange(startOfMonth, null);
```




information in
generic data structures

ACID, serialized, time aware

contains old db, new db, change


```
List newData = ...,  
Future<Map> f = conn.transactAsync(list);  
  
dbBefore = conn.db.asOf(time);  
  
possibleFuture = db.with(...);  
  
allTime = db.history();  
  
BlockingQueue<Map> queue = conn.txReportQueue();  
  
Log log = conn.log();  
Iterable<Map> it = log.txRange(startOfMonth, null);
```

ACID, serialized, time aware

```
List newData = ...;  
Future<Map> f = conn.transactAsync(list);  
  
dbBefore = conn.db.asOf(time);  time travel  
  
possibleFuture = db.with(...);  
  
allTime = db.history();  
  
BlockingQueue<Map> queue = conn.txReportQueue();  
  
Log log = conn.log();  
Iterable<Map> it = log.txRange(startOfMonth, null);
```

ACID, serialized, time aware


```
List newData = ...;  
Future<Map> f = conn.transactAsync(list);  
  
dbBefore = conn.db.asOf(time);  
  
possibleFuture = db.with(...);  
allTime = db.history();  
  
BlockingQueue<Map> queue = conn.txReportQueue();  
  
Log log = conn.log();  
Iterable<Map> it = log.txRange(startOfMonth, null);
```



one possible future

ACID, serialized, time aware

```
List newData = ...;  
Future<Map> f = conn.transactAsync(list);  
  
dbBefore = conn.db.asOf(time);  
  
possibleFuture = db.with(...);  
  
allTime = db.history();  
BlockingQueue<Map> queue = conn.txReportQueue();  
  
Log log = conn.log();  
Iterable<Map> it = log.txRange(startOfMonth, null);
```




all history, overlapped

ACID, serialized, time aware

```
List newData = ...;  
Future<Map> f = conn.transactAsync(list);  
  
dbBefore = conn.db.asOf(time);  
  
possibleFuture = db.with(...);  
allTime = db.history();  
BlockingQueue<Map> queue = conn.txReportQueue();  
  
Log log = conn.log();  
Iterable<Map> it = log.txRange(startOfMonth, null);
```

*monitor all change
from any peer*



ACID, serialized, time aware

```
List newData = ...;  
Future<Map> f = conn.transactAsync(list);  
  
dbBefore = conn.db.asOf(time);  
  
possibleFuture = db.with(...);  
  
allTime = db.history();  
  
BlockingQueue<Map> queue = conn.txReportQueue();  
  
Log log = conn.log();  
Iterable<Map> it ← log.txRange(startOfMonth, null);
```

review any
time range

complexities mitigated

lost data	managing time	eventual consistency
log analysis	test setup	DAOs
ORM	defensive copying	DTOs
inheritance	join tables	objects
structural rigidity	relationship direction	strings
model caching	logic duplication	injection attacks
app caching	imperative code	data duplication
	read transactions	isolation levels
	denormalization	

how we did it

simplicity

power

focus

courage

pragmatism

patience

complexity: tables

“People can belong to multiple clubs”

join table

person table

club table

id key in person table

person key in join table

club key in join table

id key in club table

simplicity: datoms

“People can belong to multiple clubs”

[?person :club ?club]

power

capabilities dictated at storage time	capabilities chosen at query time
row store	EAVT, partitions
kv store	AVET, unique keys
column store	AEVT
graph database	VAET
document database	all the above

agility

Docs

Rects

KV

Triples

Datoms



Rigid

Agile

leverage

KV

Docs

Rects

Triples

Datoms



Low


High

focus

"Because of the elasticity of Datomic, we were able to reduce our hosting fees by a factor of 10 when we moved off of [a popular NoSQL]."

<redacted>

courage: elastic read, ACID write



pragmatism

academic (but unfashionable) ideas

datalog

persistent data structures

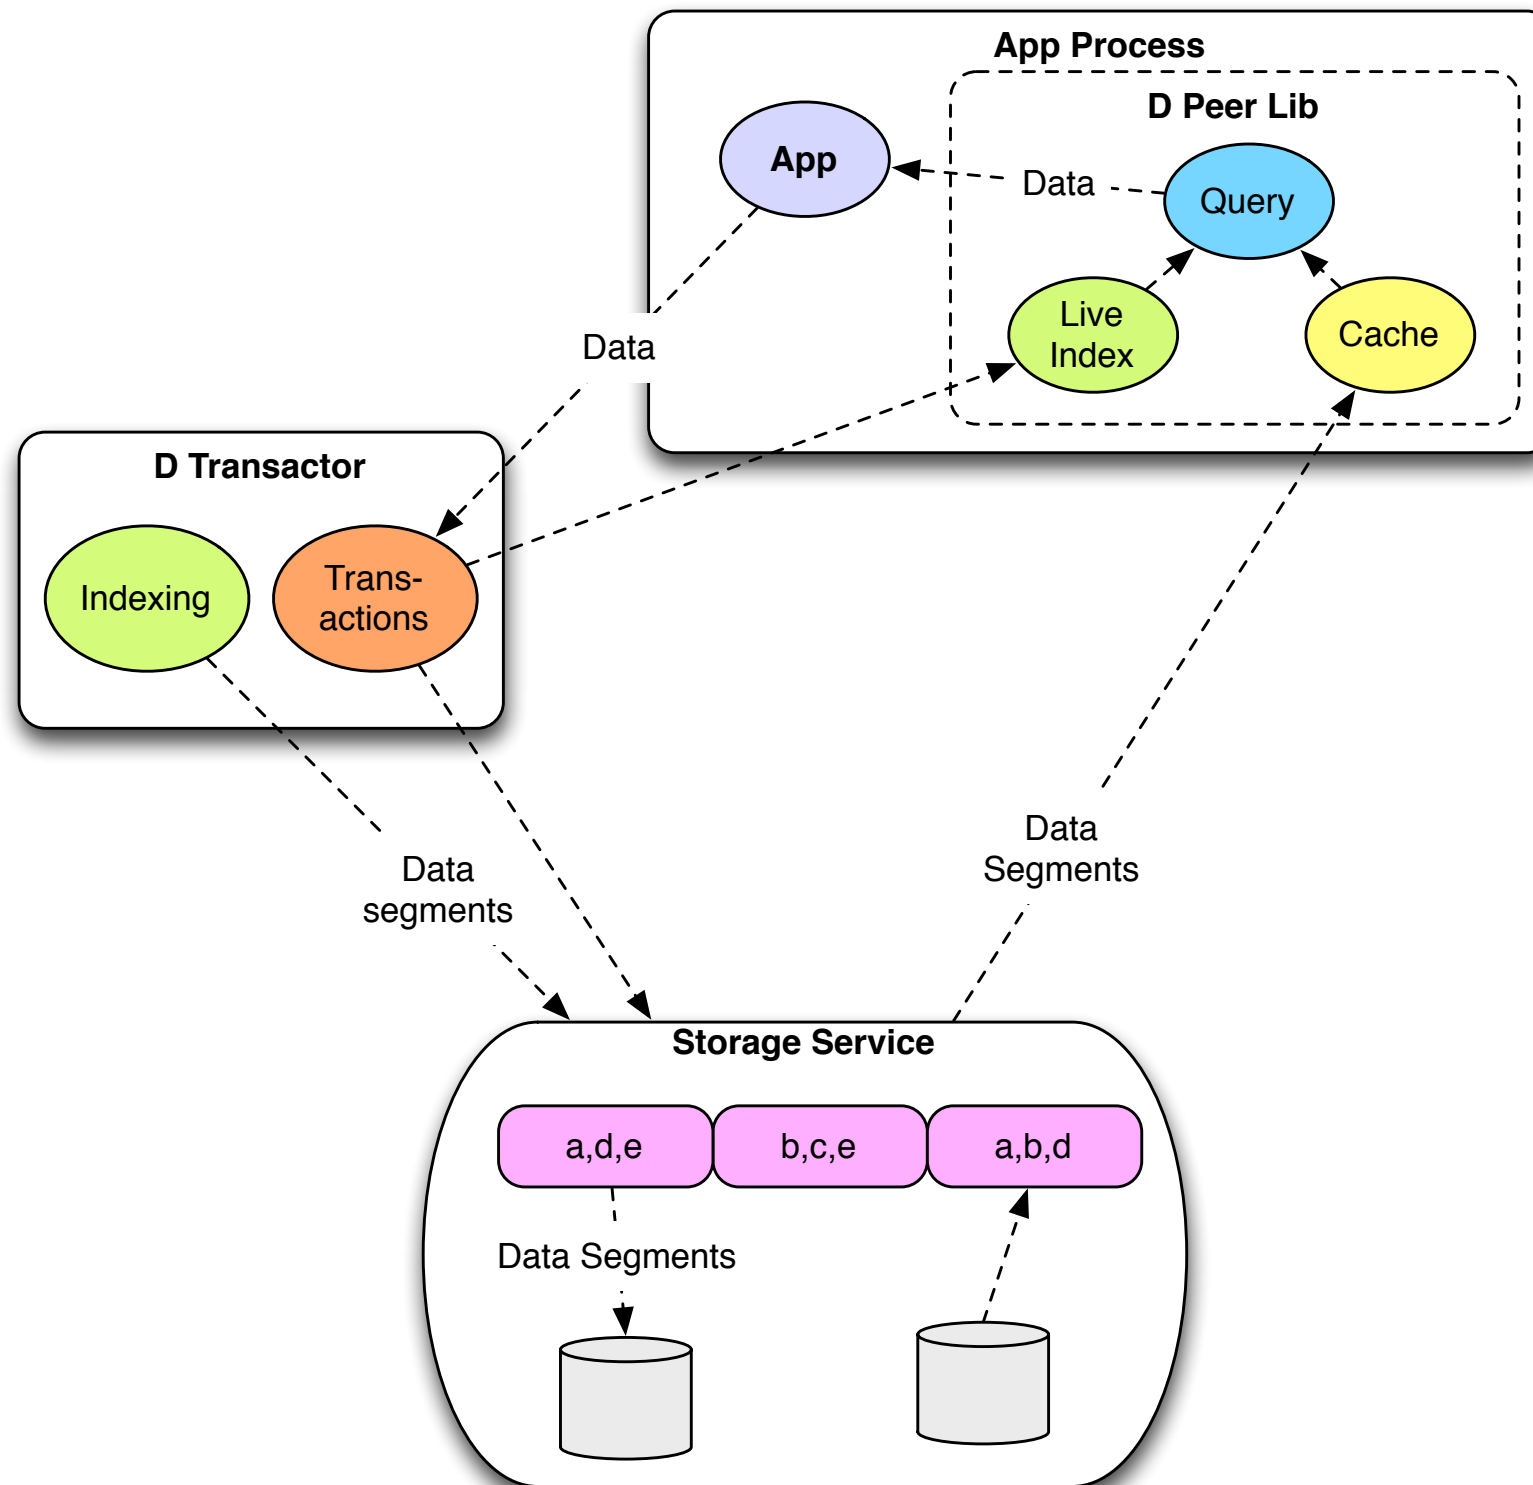
lisp

communicating sequential processes

but not these ideas:

(list omitted for troll protection)

patience



the problem with
example-based testing

example-based-tests (EBT)

```
describe Bowling, "#score" do
  it "returns 0 for all gutter game" do
    bowling = Bowling.new
    20.times { bowling.hit(0) }
    bowling.score.should eq(0)
  end
end
```

EBT

setup

```
describe Bowling, "#score" do
  it "returns 0 for all gutter game" do
    bowling = Bowling.new
    20.times { bowling.hit(0) }
    bowling.score.should eq(0)
  end
end
```


EBT


```
describe Bowling, "#score" do
  it "returns 0 for all gutter game" do
    bowling = Bowling.new
    20.times { bowling.hit(0) }
    bowling.score.should eq(0)
  end
end
```



inputs

EBT

```
describe Bowling, "#score" do
  it "returns 0 for all gutter game" do
    bowling = Bowling.new
    20.times { bowling.hit(0) }
    bowling.score.should eq(0)
  end
end
```



execution

EBT

```
describe Bowling, "#score" do
  it "returns 0 for all gutter game" do
    bowling = Bowling.new
    20.times { bowling.hit(0) }
    bowling.score.should eq(0)
  end
end
```



output

EBT

```
describe Bowling, "#score" do
  it "returns 0 for all gutter game" do
    bowling = Bowling.new
    20.times { bowling.hit(0) }
    bowling.score.should eq(0)
  end
end
```



validation

weaknesses of EBT

severely limited coverage

fragility

bespoke

deconstructing EBT

Inputs

Execution

Outputs

Validation

generative testing



loose coupling FTW

decouple

benefits

model

improve design
generate load

inputs

increase comprehensiveness by running longer

execution

test different layers with same code
only part that must change with your app

outputs

expert analysis
persist for future study

validation

test generic *properties*
run against prod data

all

functional programming
feedback loops in test development

thank you!

 **cognitect**

@stuarthalloway